

Article

Sheaf Mereology: Parts and Wholes in a Topos-Theoretic Setting

Firstname Lastname¹ , Firstname Lastname² and Firstname Lastname^{2,*}

¹ Affiliation 1; e-mail@e-mail.com

² Affiliation 2; e-mail@e-mail.com

* Correspondence: e-mail@e-mail.com; Tel.: (optional; include country code; if there are multiple corresponding authors, add author initials) +xx-xxxx-xxx-xxxx (F.L.)

Abstract

A single paragraph of about 200 words maximum. For research articles, abstracts should give a pertinent overview of the work. We strongly encourage authors to use the following style of structured abstracts, but without headings: (1) Background: place the question addressed in a broad context and highlight the purpose of the study; (2) Methods: describe briefly the main methods or treatments applied; (3) Results: summarize the article's main findings; (4) Conclusions: indicate the main conclusions or interpretations. The abstract should be an objective representation of the article, it must not contain results which are not presented and substantiated in the main text and should not exaggerate the main conclusions.

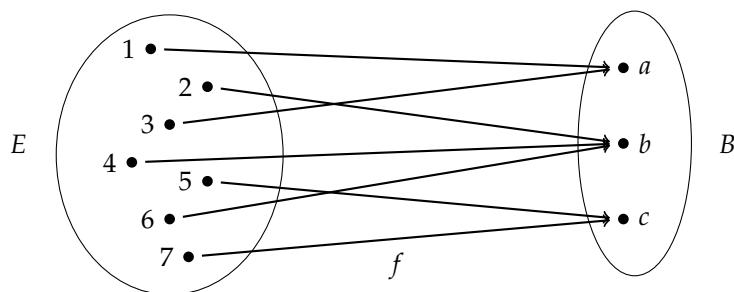
Keywords: mereology; fusions and integral wholes; sheaves; point-free topology; frames and locales; toposes; modality; merology logic; categorical logic

1. Sheaf-Theory

In this section, we cover the parts of sheaf theory that we will utilize in the rest of the paper. Readers familiar with sheaf theory can skip this section.

1.1. Fibers

Suppose we have a map (function) $f : E \rightarrow B$ that looks something like this:



Received:

Revised:

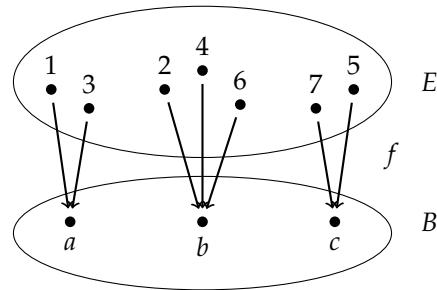
Accepted:

Published:

Citation: Lastname, F.; Lastname, F.; Lastname, F. Title. *Philosophies* **2025**, *1*, 0. <https://doi.org/>

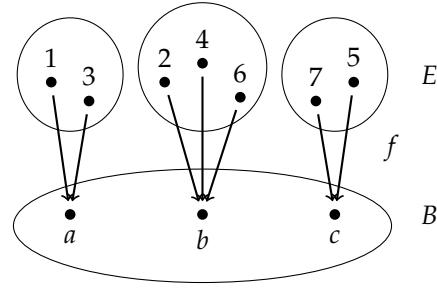
Copyright: © 2025 by the authors. Submitted to *Philosophies* for possible open access publication under the terms and conditions of the Creative Commons Attribution (CC BY) license (<https://creativecommons.org/licenses/by/4.0/>).

It is sometimes convenient to turn the diagram sideways and group together points in the domain that get sent to the same target, like so:



21

That makes the pre-images very easy to see. For any point in B , its pre-image is just the group of points sitting “over” it:

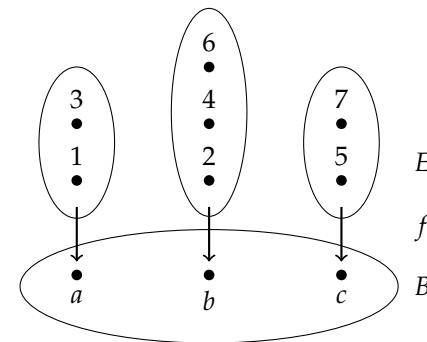


22

23

24

If we stack the points in each pre-image vertically, one on top of the other, we can then think of each pre-image as a kind of “stalk” growing over its base point:



25

26

27

This gives rise to the idea of the “fibers” of a map. The fibers of a map are just its pre-images. For instance, the fiber over b is $\{2, 4, 6\}$.

28

29

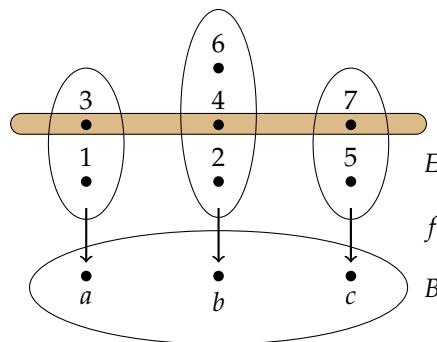
Definition 1 (Fibers). Given a map $f : E \rightarrow B$ and a point $y \in B$, the fiber over y is its pre-image $f^{-1}(y) = \{x \mid f(x) = y\}$. B is called the “base space” (or “base” for short) of f , while the point y is called the “base point” (or “base” for short) of the fiber.

30

31

32

We can select a cross-section of one or more fibers by selecting a point from each of the fibers in question. For instance, we can take 3, 4, and 7 as a cross-section of the fibers $f^{-1}(a)$, $f^{-1}(b)$, and $f^{-1}(c)$:



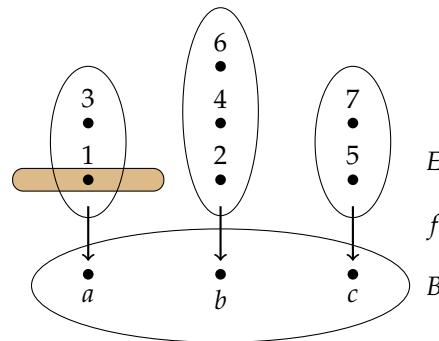
33

34

35

36

We can also take cross-sections local to only some of the fibers. For instance, we can take 1 as a cross-section just of $f^{-1}(a)$:



Definition 2 (Sections). Given a map $f : E \rightarrow B$ and a subset $C \subseteq B$ (i.e., a selection of base points in B), a section of f (over C) is a choice of one element from each fiber over each base $x \in C$.

Remark 1. Since each point in a fiber amounts to a section over the fiber's base, the elements of a fiber are often just called the "sections" of the fiber.

1.2. Spaces

In the above examples, the base B was a set. We often want to consider bases that have more structure, e.g., bases that have spatial structure.

In traditional topology, spaces are built out of the points of the space. Given a set of points, a topology on that set specifies which points belong in which regions of the space.

Definition 3 (Topology). Let X be a non-empty set, thought of as the "points" of the space. A topology on X is a collection T of subsets of X , thought of as the "regions" of the space (called the "open sets" or just the "opens" of T), that satisfy the following conditions:

(T1) The empty set and the whole set are open:

$$\emptyset \in T, X \in T.$$

(T2) Arbitrary unions of opens are open:

$$\text{if } \{U_i\}_{i \in I} \subseteq T, \text{ then } \bigcup_{i \in I} U_i \in T.$$

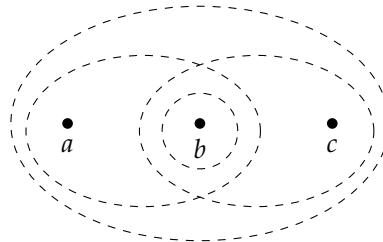
(T3) Finite intersections of opens are open:

$$\text{if } U_1, \dots, U_n \in T, \text{ then } \bigcap_{i=1}^n U_i \in T.$$

These conditions encode the way that spatial regions are put together. For instance, it ensures that if two regions overlap, then their overlapping area is a region too (and indeed, that's what it means for regions to *overlap*: there's a region of space they have in common).

Remark 2. The regions of a topology, ordered by inclusion, form a complete lattice. Since the topology includes arbitrary unions, the join of this lattice is set union, but since the topology includes only finite intersections, the meet of this lattice is the interior of set intersection.

Example 1. Let $X = \{a, b, c\}$. One possible topology is: $T = \{\emptyset, \{b\}, \{a, b\}, \{b, c\}, \{a, b, c\}\}$. If we draw dashed circles around each of the opens (regions), ignoring the empty set, we get:



63

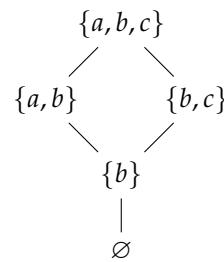
There are two regions $\{a, b\}$ and $\{b, c\}$ that overlap at b (so $\{b\}$ is a region in T too). There is also the full region $\{a, b, c\}$, which is the union of the smaller regions.

64

65

We can draw T as a Hasse diagram, which shows that the regions form a lattice:

66



67

The lattice structure suggests that much of what is important about a space is not so much its points, but rather its opens/regions. This leads to the idea that topology-like reasoning can be done without the points. So, we can generalize: take a topology, and drop the points. That leaves just the opens/regions, which we call a frame (or locale).

68

69

70

71

Definition 4 (Frames/locales). A frame (synonymously, a locale) \mathbb{L} is a partially ordered set L (we call its elements “opens” or “regions”) that satisfies the following conditions:

72

73

(L1) L is a complete lattice:

74

- Every subset $S \subseteq L$ has a join, denoted $\vee S$.
- Every finite subset $F \subseteq L$ has a meet, denoted $\wedge F$.

75

76

(L2) Finite meets distribute over arbitrary joins:

$$a \wedge \bigvee_{i \in I} b_i = \bigvee_{i \in I} (a \wedge b_i), \text{ for all } a \in L \text{ and all families } \{b_i\}_{i \in I} \subseteq L.$$

Define $V \preceq U$ (read “ V is included in U ”) by $a = a \wedge b$.

77

Remark 3. The fact that $V \preceq U$ is equivalent to $a = a \wedge b$ means we can deal with the opens of a frame algebraically (via \wedge and \vee operations), or order-theoretically (via the \preceq relation), whichever is more convenient.

78

79

80

Remark 4. The category of locales is defined as the dual/opposite of the category of frames (see ?? below), and so frames and locales are quite literally the very same objects. In practice, frames are often used for algebraic purposes, and locales are used for (generalized) spatial purposes. Here, we will have no reason to distinguish these two roles, and so we will use the names “frame” and “locale” interchangeably.

81

82

83

84

85

1.3. Presentations of locales

86

Locales have presentations much like groups and other algebraic structures. To give the presentation of a locale, specify a set of generators and relations.

87

88

Definition 5 (Presentations). A presentation $\langle G, R \rangle$ of a locale \mathbb{L} is comprised of:

89

(P1) A set of generators $G = \{U_k, U_m, \dots\}$. 90

(P2) A set of relations $R \subseteq G \times G$ on those generators. 91

The locale \mathbb{L} presented by $\langle G, R \rangle$ is the smallest one freely generated from G which satisfies R . 92

Remark 5. Every locale has a presentation, and a locale can have multiple presentations. 93

To calculate the locale that corresponds to a presentation, start with the generators, then take all finite meets and all arbitrary joins that satisfy R (and of course L1 and L2). 94
95

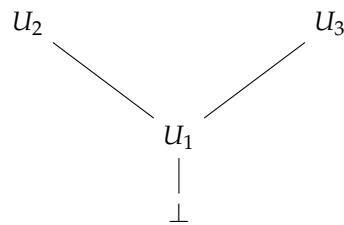
Example 2. Let a locale \mathbb{L} be given by the presentation $\langle G, R \rangle$ where: 96

- $G = \{\perp, U_1, U_2, U_3\}$. 97
- $R = \{\perp \preceq U_1, U_1 \preceq U_2, U_1 \preceq U_3\}$. 98

There are four generators (\perp, U_1, U_2 , and U_3), and \perp is below U_1 while U_1 is a sub-region of U_2 and U_3 . Since U_1 is a sub-region of both U_2 and U_3 , U_1 is their meet: 99
100

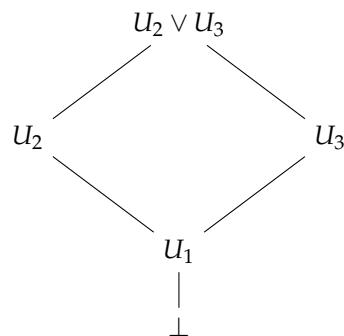
- $U_1 = U_2 \wedge U_3$. 101

At this point, we have generated this much of the locale: 102



103

R says nothing to constrain joins, so we need to join everything we can. In this case, we need to join U_2 and U_3 : 104
105



106

There are no further joins or meets that aren't already represented in the picture. For instance, all further non-trivial meets are already accounted for: 107
108

- $U_1 \wedge \perp = \perp$. 109
- $U_2 \wedge U_1 = U_1$ and $U_3 \wedge U_1 = U_1$. 110
- $U_2 \wedge \perp = \perp$ and $U_3 \wedge \perp = \perp$. 111
- $(U_2 \vee U_3) \wedge U_2 = U_2$ and $(U_2 \vee U_3) \wedge U_3 = U_3$. 112
- $(U_2 \vee U_3) \wedge U_1 = U_1$. 113
- $(U_2 \vee U_3) \wedge \perp = \perp$. 114

Similarly, all other non-trivial joins are also already accounted for: 115

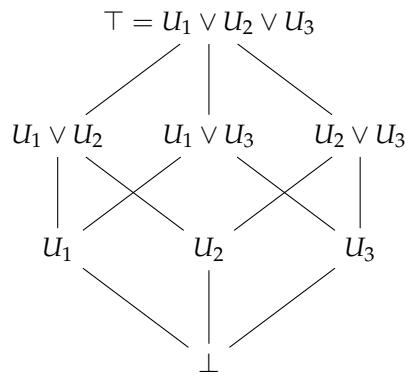
- $\perp \vee U_1 = U_1$. 116
- $\perp \vee U_2 = U_2$ and $\perp \vee U_3 = U_3$. 117
- $\perp \vee (U_2 \vee U_3) = U_2 \vee U_3$. 118

- $U_1 \vee U_2 = U_2$ and $U_1 \vee U_3 = U_3$. 119
- $U_2 \vee (U_2 \vee U_3) = U_2 \vee U_3$ and $U_2 \vee (U_3 \vee U_3) = U_2 \vee U_3$. 120

Example 3. Let $\mathbb{L} = \langle G, R \rangle$ be given by:

- $G = \{U_1, U_2, U_3\}$. 121
- $R = \emptyset$. 122

We have three generators (U_1 , U_2 , and U_3), and there are no relations restricting how those generators are related. Thus, the locale that is freely generated from this presentation is isomorphic to the power set of three elements:



127

A presentation provides the most “minimal” information from which the rest of the locale is generated. It might be tempting to think that each generator is an atomic region, but that is not quite right. Some generators are reducible to others.

128

129

130

Definition 6 (Meet-irreducibility). Given a presentation $\mathbb{L} = \langle G, R \rangle$, a region $U \in G$ is meet-reducible (“reducible” for short) if it is the non-trivial meet of other regions:

131

132

$$\exists W, V \in \mathbb{L}, U \neq W, U \neq V, \text{ and } U = W \wedge V.$$

U is meet-irreducible (“irreducible” for short) if it is not meet-reducible, i.e.:

133

$$\forall W, V \in \mathbb{L}, U = W \wedge V \implies (V = U \text{ or } W = U).$$

Intuitively, a generator is reducible if it can be expressed as the meet of strictly larger regions, which occurs exactly when it is their overlap.

134

135

Example 4. Take the locale from Example 2. U_1 is the overlap of U_2 and U_3 , and U_2 and U_3 are strictly larger regions than U_1 , so U_1 is reducible.

136

137

By contrast, U_2 and U_3 are irreducible, because they cannot be expressed as the meet of two strictly larger regions. Similarly, \perp is irreducible, because it is not the meet of two strictly larger regions either (it is the meet of only one strictly larger region, namely U_1).

138

139

140

We can see the minimal irreducible generators of a locale as its atomic regions.

141

Definition 7 (Atomic regions). Given a presentation $\mathbb{L} = \langle G, R \rangle$, define the atomic regions of \mathbb{L} , denoted $\text{Atoms}(\mathbb{L})$, as the minimal irreducible generators, i.e. those generators $g \in G$ that satisfy the following two conditions:

142

143

144

(A1) Meet-irreducibility. g is meet-irreducible.

145

(A2) Minimality. There is no strictly smaller meet-irreducible h with $h \preceq g$.

146

use “primitive” or “fundamental” instead of “atomic”? An atomic thing can’t be broken down, but a meet-irreducible element does have a part!

147

1.4. Presheaves

148

Above we considered the fibers of a map $f : E \rightarrow B$, where E and B were sets. We can also consider fibers over locales, where the fibers respect the locale’s structure. This is called a presheaf. A presheaf is an assignment of data to each of a locale’s regions that is “stable under restriction,” i.e., that respects zooming in and out.

149

150

151

152

Definition 8 (Presheaf). Let \mathbb{L} be a locale, and let $Morphs(\mathbb{L})$ be $\{\langle A, B \rangle \mid A \preceq B \in \mathbb{L}\}$. A presheaf on \mathbb{L} is a pair $\langle F, \{\rho_A^B\}_{\langle A, B \rangle \in Morphs(\mathbb{L})} \rangle$, where:

153

154

- F assigns to each region $U \in \mathbb{L}$ some data $F(U)$.
- $\{\rho_A^B\}_{\langle A, B \rangle \in Morphs(\mathbb{L})}$ is a family of maps $\rho_A^B : F(B) \rightarrow F(A)$ (called “restriction maps”), each of which specifies how to restrict the data over $F(B)$ down to the data over $F(A)$.

155

156

157

All together, $\langle F, \{\rho_A^B\}_{\langle A, B \rangle \in Morphs(\mathbb{L})} \rangle$ must satisfy the following conditions:

158

(R1) Restrictions preserve identity:

$$\rho_U^U = id_U \text{ (the identity on } U\text{), for every } U \in \mathbb{L}.$$

(R2) Restrictions compose:

$$\text{If } A \preceq B \text{ and } B \preceq C, \text{ then } \rho_A^C = \rho_A^B \circ \rho_B^C.$$

Since F assigns data $F(U)$ to each region $U \in \mathbb{L}$, we can think of the $F(U)$ s as the “fibers” over \mathbb{L} , and the restriction maps as “zoom in” maps that go from bigger fibers down to smaller fibers.

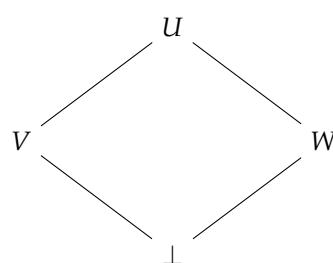
159

160

161

Example 5. Let \mathbb{L} be a locale $\{\perp, W, V, U\}$ with the following structure:

162



163

Next, let’s define a presheaf F as follows:

164

- $F(U) = \{a, b, c, d\}$, $F(V) = \{a, b\}$, $F(W) = \{c, d\}$, $F(\perp) = \{*\}$.
- Define ρ_V^U as the projection (send a to a , b to b , and the rest can go anywhere), and similarly for ρ_W^U . Let ρ_{\perp}^U , ρ_{\perp}^V , and ρ_{\perp}^W send their data to $\{*\}$, and let the rest be identities.

165

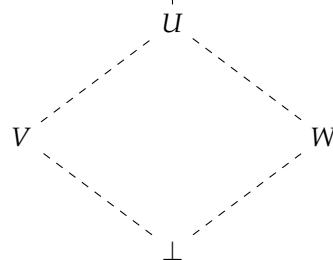
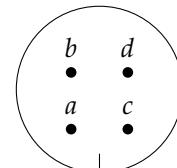
166

167

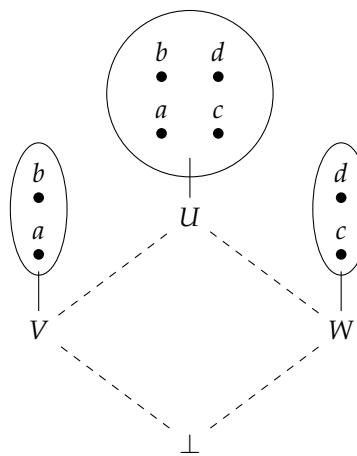
We can see F ’s assignments as fibers over \mathbb{L} by drawing them over the regions they are assigned to. For instance, over U we have $F(U)$, i.e., $\{a, b, c, d\}$:

168

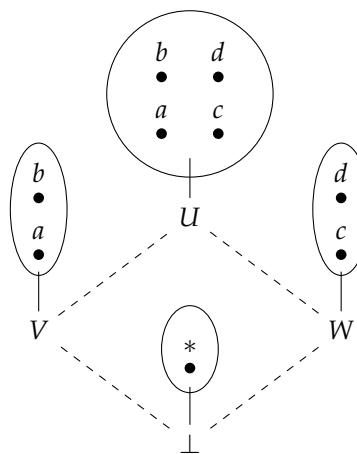
169



Similarly, over V and W , we have $F(V) = \{a, b\}$ and $F(W) = \{c, d\}$:



Finally, over \perp , we have a singleton set:

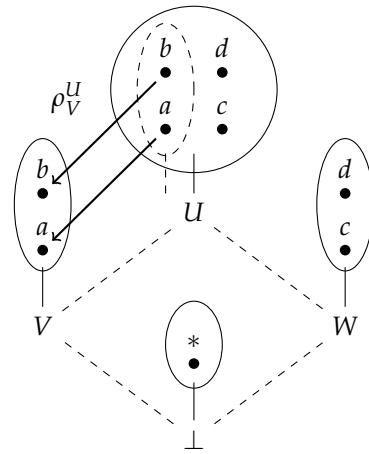


The restriction maps show how to “zoom in” on the parts (sub-fibers) of any given fiber. For instance, we can see that the fiber over V is contained in the fiber over U . The restriction map just projects that sub-fiber out, thereby showing us how to “zoom in” on that sub-fiber:

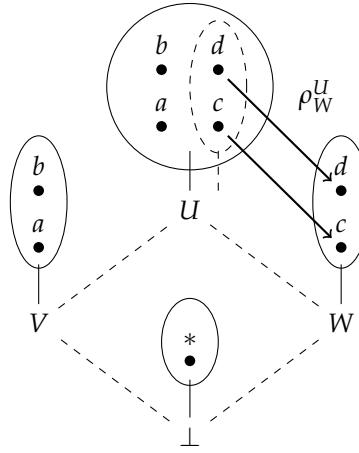
175

176

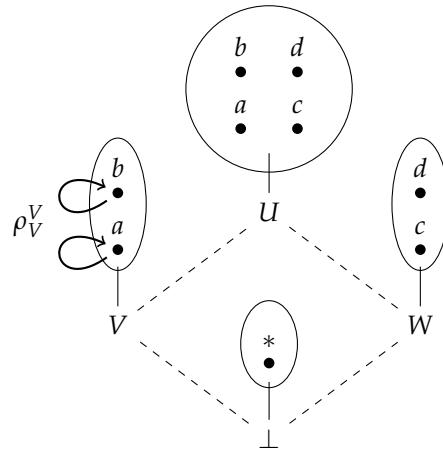
177



178

It's similar for the fiber over W:

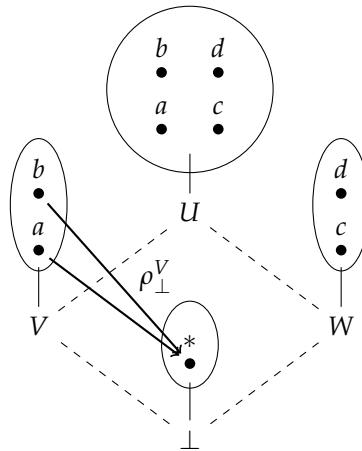
179

Restricting a fiber to itself is just the identity on the fiber:

180

The other restriction maps restrict down to the singleton set. For instance:

183

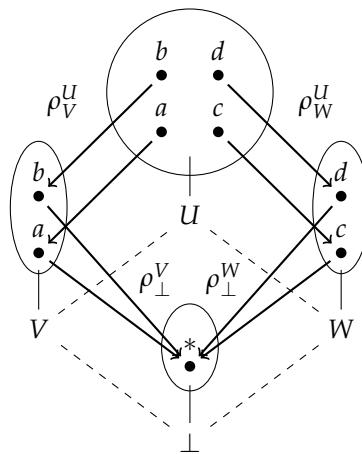


184

All of this makes it clear that the structure of the presheaf data that sits in the fibers over \mathbb{L} mimics (respects) the structure of the base locale:

185

186



187

Remark 6. In line with Remark 1, the elements of each fiber $F(U)$ are usually just called the “sections” of $F(U)$. For instance, c is a section of $F(W)$, as is d .

188

189

1.5. Sheaves

190

The definition of a presheaf requires only that the data be stable under restriction (zooming in on a region). It does not require that the data fit together across different regions (fibers).

191

192

A sheaf is a presheaf with an added gluing condition: whenever you have compatible data on overlapping fibers, there must be a unique way to glue it together into data over the union. In other words, the data in the fibers must agree on overlap and combine coherently.

193

194

195

To get at this idea, let’s first define a cover. A cover of a region U is a selection of sub-regions that covers U in its entirety. The chosen sub-regions don’t leave any part of U exposed.

196

197

198

Definition 9 (Cover). Let \mathbb{L} be a topology or a locale, and let U be a region of \mathbb{L} . A cover of U is a family $\{U_i\}_{i \in I} \subseteq \mathbb{L}$ such that:

199

200

$$U = \bigvee_{i \in I} \{U_i\}.$$

In other words, a cover of U is a family of regions that join together to form U .

201

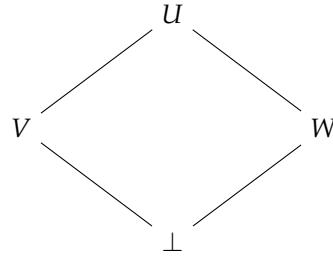
Example 6. Take the topology from Example 1: $T = \{\emptyset, \{b\}, \{a, b\}, \{b, c\}, \{a, b, c\}\}$. A cover of $\{a, b, c\}$ is $\{a, b\}$ and $\{b, c\}$, because altogether, $\{a, b\}$ and $\{b, c\}$ cover all of the points in $\{a, b, c\}$.

202

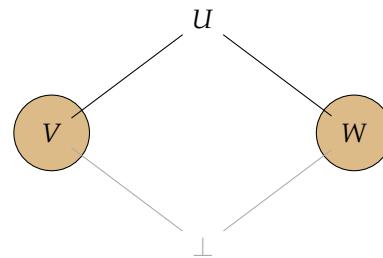
203

Another cover of $\{a, b, c\}$ is $\{\{a, b\}, \{b, c\}, \{b\}\}$. Although $\{b\}$ is redundant here, this choice of sub-regions still entirely covers $\{a, b, c\}$ as required.

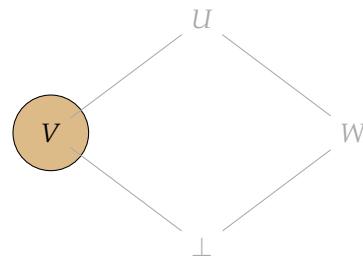
Example 7. In the context of frames, where there are no points, a cover of U is just a selection of sub-regions of U that together join together to form U . Take the locale from Example 5:



A cover of U is $\{V, W\}$, since $U = \bigvee \{V, W\}$:



A cover of V is just $\{V\}$:



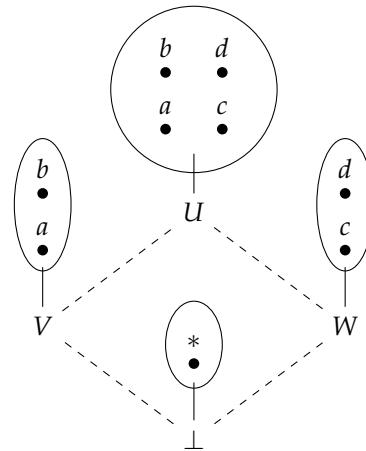
Remark 7. A cover over the least element of a locale (or a topology) is empty (the empty set), because there are no regions (or points) to cover.

Given a presheaf F over a locale \mathbb{L} , if we have a cover $\{U_i\}_{i \in I}$ of some portion of \mathbb{L} , there is a corresponding family of fibers $\{F(U_i)\}_{i \in I}$ over that cover. We can pick one section (i.e., one element) from each such fiber to get a slice of elements that spans all of the fibers over that cover. Let us call such a choice a selection of patch candidates.

Definition 10 (Patch candidates). Given a presheaf F and a cover $\{U_i\}_{i \in I}$ with a corresponding family of fibers $\{F(U_i)\}_{i \in I}$, a selection of patch candidates $\{s_i\}_{i \in I}$ is a choice of one section s_i from each $F(U_i)$:

$$\{s_i\}_{i \in I} = \{s_i \mid s_i \in F(U_i) \text{ for each } F(U_i) \in \{F(U_i)\}_{i \in I}\}.$$

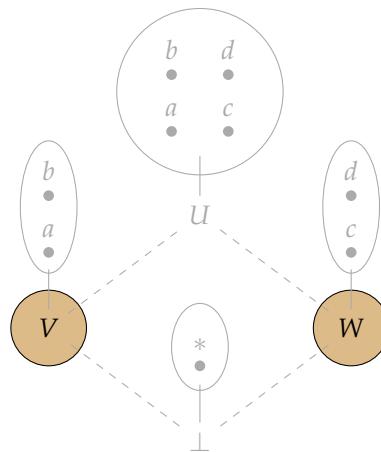
Example 8. Take the presheaf from Example 5:



223

Let $\{V, W\}$ be the cover of interest:

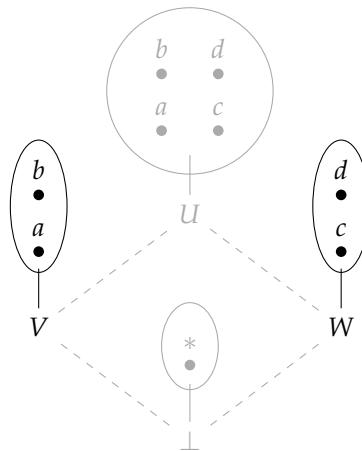
224



225

Over this cover, we have a corresponding family of fibers:

226

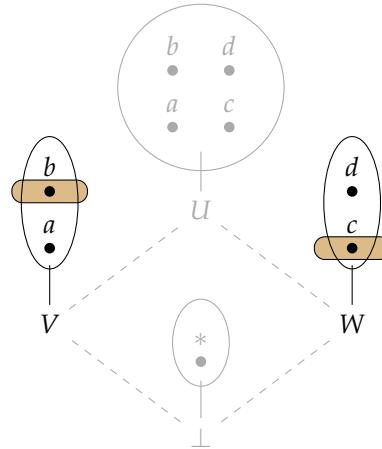


227

A selection of patch candidates is a choice of one section (element) from each fiber. For instance, we might pick b from $F(V)$ and c from $F(W)$:

228

229



230

Similarly, we might pick $\{a, d\}$, $\{b, d\}$, or $\{a, c\}$, each of which is a valid selection of patch candidates.

231

232

Remark 8. Consider the empty cover. Since there are no sub-regions below the least element of a locale, there are no patch candidates we could choose for the empty cover either. Hence, any selection of patch candidates for the empty cover is \emptyset .

233

234

235

A selection of patch candidates might fit together, or they might not. We say they are compatible if they fit together, i.e., if they agree on overlaps. To check this, take any pair of patch candidates, and check if they restrict to the same data on their overlap.

236

237

238

Definition 11 (Compatible patch candidates). Given two fibers $F(U_i)$ and $F(U_j)$ and a patch candidate from each, $s_i \in F(U_i)$ and $s_j \in F(U_j)$, s_i and s_j are compatible if they restrict to the same data on their overlap $U_i \wedge U_j$:

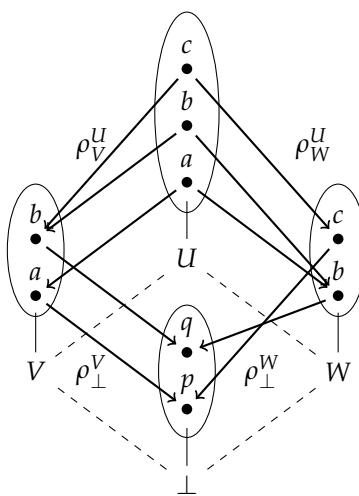
$$\rho_{U_i \wedge U_j}^{U_i}(s_i) = \rho_{U_i \wedge U_j}^{U_j}(s_j).$$

A selection of patch candidates $\{s_i\}_{i \in I}$ is compatible if all of its members are pair-wise compatible.

239

Example 9. Consider the following presheaf F :

240

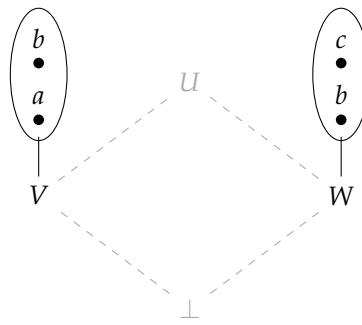


241

Take the cover $\{V, W\}$ and its corresponding fibers:

242

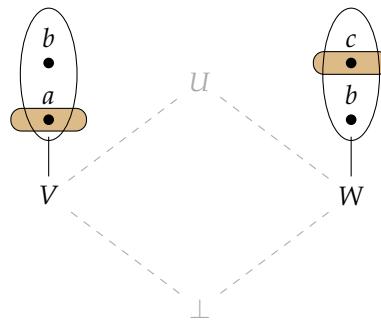
243



243

Suppose we pick $\{a, c\}$ for patch candidates:

244



245

Is this selection compatible? We have to check if they agree on their overlap. The overlap $V \wedge W$ is \perp . Where does ρ_{\perp}^V send our chosen patch candidate a ? It sends it to p , since $\rho_{\perp}^V(a) = p$. Where does ρ_{\perp}^W send our other chosen patch candidate b ? It also sends it to p , since $\rho_{\perp}^W(c) = p$. On the overlap \perp then, $\rho_{\perp}^V(a) = \rho_{\perp}^W(c)$, so a and c are compatible. This is easy to see in the diagram, since a and c both get sent to the same place:

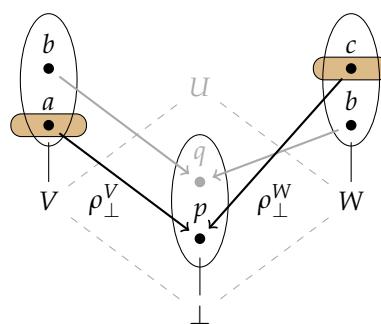
246

247

248

249

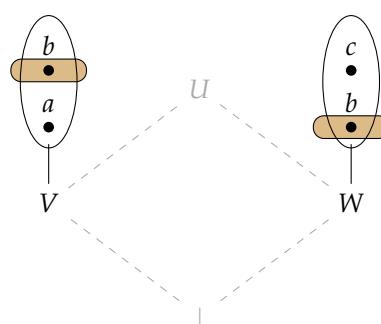
250



251

Now suppose we pick $\{b, b\}$ for patch candidates:

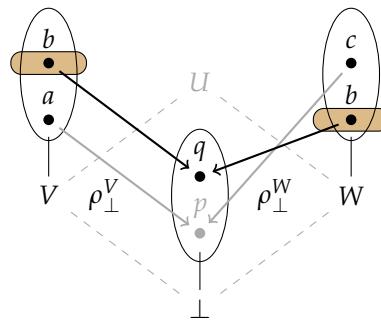
252



253

These are also compatible. They agree on their overlap (both restrict to q):

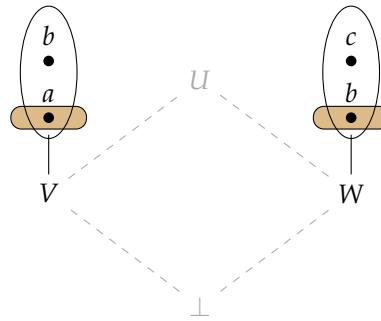
254



255

Finally, suppose we pick $\{a, b\}$ for patch candidates:

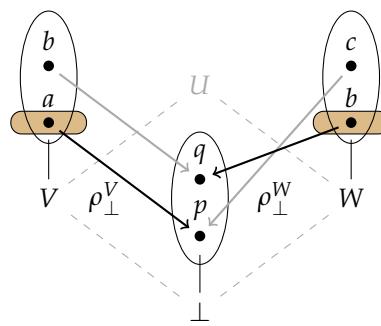
256



257

These are not compatible. They do not agree on their overlap:

258



259

Remark 9. Consider the empty cover. Since any selection of patch candidates for the empty cover is empty, compatibility is satisfied vacuously.

260

261

As an analogy, if you ask your class to turn off all cell phones but nobody brought a cell phone to class, then your request is satisfied vacuously: there is simply nothing that needs to be done to make it happen. It's similar with the empty cover: since there are no patch candidates to check, compatibility is achieved vacuously.

262

263

264

265

If a selection of patch candidates s_i, \dots, s_k across a cover of U is compatible, we say those patches glue together if if there's a section s in $F(U)$ that restricts down to exactly those patches.

266

267

268

Definition 12 (Gluing). Given a presheaf F and a selection of compatible patch candidates $\{s_i\}_{i \in I}$ for a cover $\{U_i\}_{i \in I}$, $\{s_i\}_{i \in I}$ glue together only if there is a section $s \in F(U)$ that restricts down to s_i on each fiber $F(U_i)$ of the cover, i.e., only if s is such that:

269

270

271

$$\rho_{U_i}^U(s) = s_i, \text{ for each } i \in I.$$

As a matter of terminology, if a section $s \in F(U)$ is glued from patches $\{s_i\}_{i \in I}$, we say that s is a global section of the cover, and each s_i is a local section of the cover. We may also say variously that

272

273

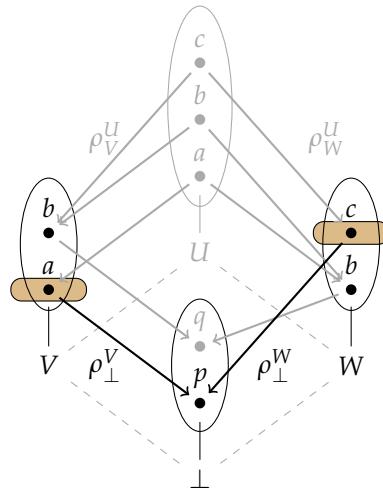
s is glued from those patches, that s is composed of those patches, that those patches compose s, or that gluing those patches yields s.

274
275

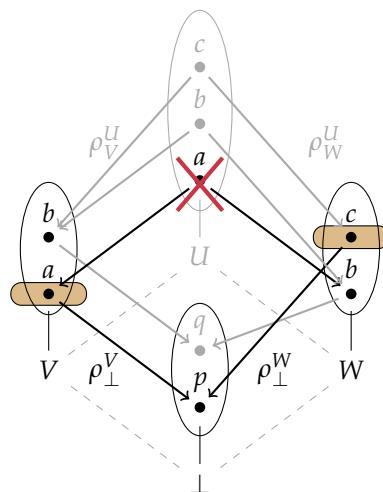
A selection of patches $\{s_i\}_{i \in I}$ glues uniquely if there is one and only one such section $s \in F(U)$ that is glued from them.

276
277

Example 10. Take the presheaf from Example 9, and consider the cover $\{V, W\}$ again. Take the the selection of patches $\{a, c\}$, which are compatible because they agree on overlap:

278
279280
281
282
283

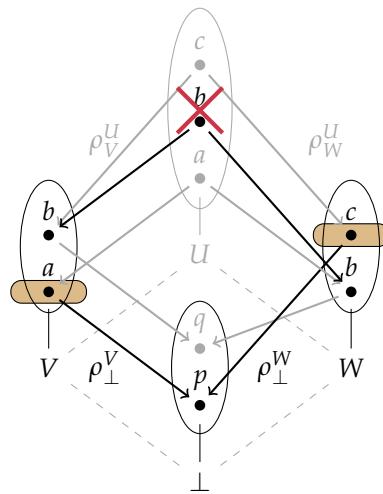
Even though a and c are compatible, they do not glue together, because there is no section in $F(U)$ that restricts down to them. Consider $a \in F(U)$ first. It restricts to $a \in F(V)$ on the left, but it does not restrict to $c \in F(W)$ on the right:



284

As for $b \in F(U)$, it restricts to neither $a \in F(V)$ on the left nor $c \in F(W)$ on the right:

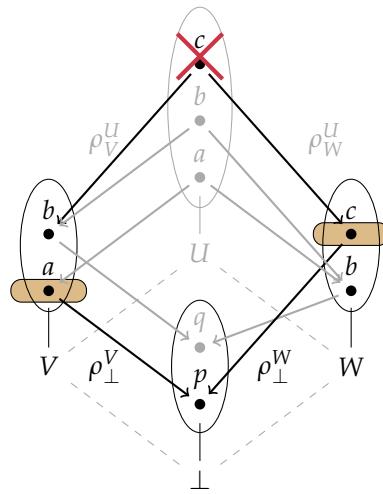
285



286

Finally, $c \in F(U)$ restricts to $c \in F(W)$ on the right, but not to $a \in F(V)$ on the left:

287



288

Thus, none of a , b , or c in $F(U)$ are glued from $\{a, c\}$, because none of them decompose into a on the left and c on the right.

289

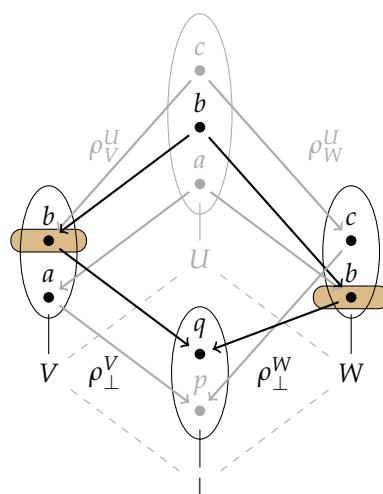
290

Now suppose we pick $\{b, b\}$ for patch candidates. These do glue together (trivially), because there is a section in $F(U)$ (namely $b \in F(U)$) that restricts down to $b \in F(V)$ on the left and $b \in F(W)$ on the right:

291

292

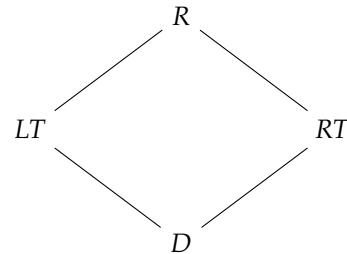
293



294

Example 11. Consider an example that glues together behaviors. Imagine a toy robot that looks something like a small tank: it has tracks on the left and right sides, and the two tracks are connected by a single drive controller. The controller either drives at a constant speed, or it sits idle. When it drives, it turns both tracks at the same speed.

Let's represent the robot as a locale. Let LT and RT be the left and right track assemblies respectively, let D be the drive controller that is shared by LT and RT , and let R be the whole robot (the join of LT and RT). As a picture:



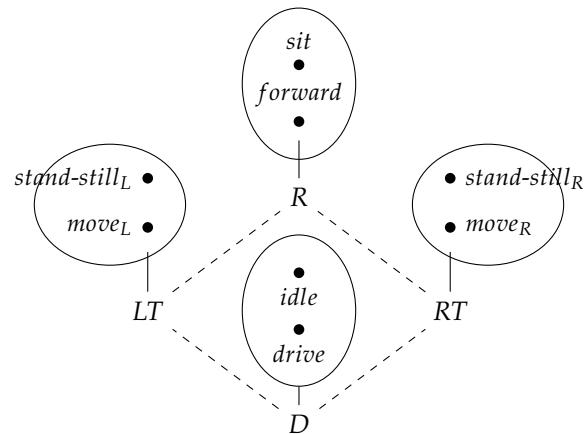
295
296
297
298
299
300

302
303
304
305
306
307

For a presheaf, let's assign to each region the behaviors that are locally observable at that region:

- The drive controller D can either drive (turn) or sit idle.
- The left track assembly can each either $move_L$ or stand-still $stand-still_L$.
- The right track assembly can also either $move_R$ or stand-still $stand-still_R$.
- The entire robot can either move forward or sit stationary.

In a picture:



309
310
311
312
313

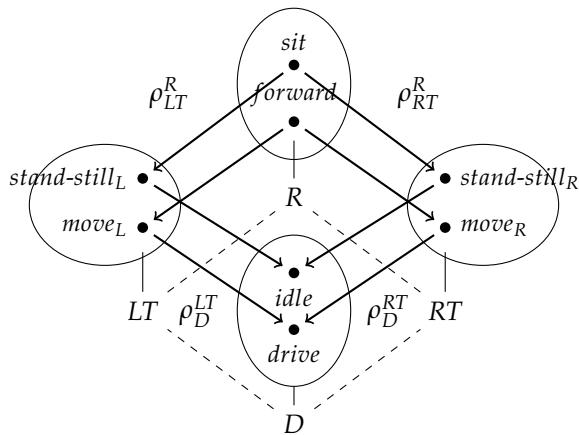
For the restriction maps, let's say that they restrict the observable behavior of a larger region to the observable behavior of the smaller region. For instance, if you are observing the whole robot moving forward (forward), and you then "zoom in" on the left track assembly, you'll see those tracks rotating ($move_L$).

- $\rho_{LT}^R(\text{sit}) = \text{stand-still}_L, \rho_{LT}^R(\text{forward}) = \text{move}_L$.
- $\rho_{RT}^R(\text{sit}) = \text{stand-still}_R, \rho_{RT}^R(\text{forward}) = \text{move}_R$.
- $\rho_D^{LT}(\text{stand-still}_L) = \text{idle}, \rho_D^{LT}(\text{move}_L) = \text{drive}$.
- $\rho_D^{RT}(\text{stand-still}_R) = \text{idle}, \rho_D^{RT}(\text{move}_R) = \text{drive}$.

In a picture:

314
315
316
317

318



319

Now take the cover $\{LT, RT\}$ of R . The patch candidates $\{move_L, move_R\}$ are compatible, because they agree on overlap (they both restrict down to drive). But they also glue uniquely, yielding forward. In other words, the robot's forward motion is patched together precisely from the two pieces of its cover, namely the left tracks rotating ($move_L$) and the right tracks rotating ($move_R$). 320
321
322
323
324

Similarly, the Robot's sitting still (sit) behavior is also glued from the two pieces of its cover, namely the left track assembly standing still ($stand-still_L$) and the right track assembly standing still ($stand-still_R$). 325
326
327

Thus, there are two global sections of R 's behavior: moving forwards (patched together from its left and right motions), or standing still (patched together from its left and right lack of motion). 328
329

We can now state what it is to be a sheaf. A sheaf is a presheaf that satisfies a special gluing condition: namely, that for every cover, every compatible selection of patch candidates glues together uniquely. 330
331
332

Definition 13 (Sheaf). A presheaf F is a sheaf iff it satisfies the following condition (called “the gluing condition”): 333
334

(G0) For every cover $\{U_i\}_{i \in I}$ of a region U and every selection of patch candidates $\{s_i\}_{i \in I}$ for that cover, if $\{s_i\}_{i \in I}$ are compatible, then there exists a unique gluing $s \in F(U)$ of $\{s_i\}_{i \in I}$. 335
336

Remark 10. There is a subtlety regarding what sheaves look like over the least element of a locale. Note that the gluing condition is formulated as an implication. That is to say, it says that, for every cross-section of patch candidates, if that cross-section can glue, then it glues in exactly one way. 337
338
339

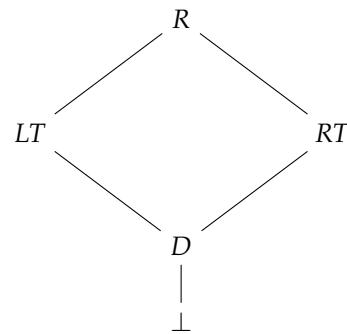
Next, consider the fact that the cover over the least region of a locale is an empty cover. Since there are no patch candidates that need to be checked for compatibility, there is nothing that needs to be done to get a “selection of gluable patch candidates.” Hence, the antecedent of the gluing condition is satisfied vacuously over the least element of the locale. 340
341
342
343

But since the empty cover satisfies the antecedent of the gluing condition vacuously, it follows that if a presheaf is to qualify as a sheaf, it must ensure that the consequent is satisfied over the empty cover as well. In other words, it must assign a unique glued section (a singleton set) to the least region of the locale. So, even though a presheaf may assign a larger set of data to the least element of a locale, a sheaf always assigns a singleton to that region. 344
345
346
347
348

Example 12. The presheaf from Example 9 fails to be sheaf, because as we saw in Example 10, there is a compatible selection of patch candidates (namely, $\{a, c\}$) which fails to glue. To be a sheaf, every compatible selection of patch candidates must glue. 349
350
351

Example 13. The presheaf from Example 11 fails to be a sheaf, because it does not assign a singleton to the lowest region of the underlying locale. In that example D is the lowest region of the locale, and $F(D) = \{p, q\}$, a set containing two elements. Hence, F fails to be a sheaf.

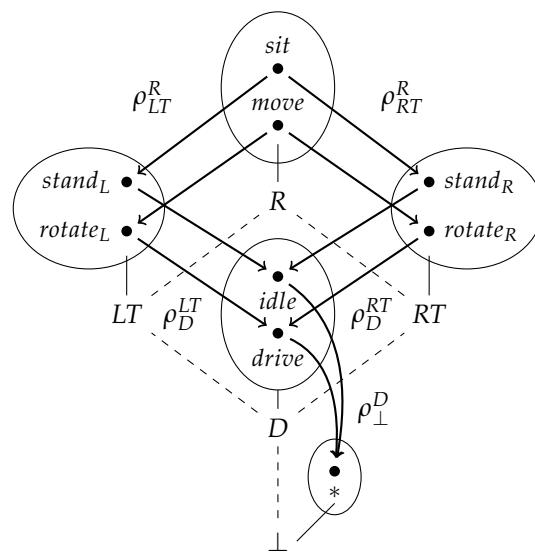
However, suppose we add a distinct bottom element to the locale:

352
353
354
355

356

If we assign a singleton set to \perp (so that $F(\perp) = \{\ast\}$, say), then F looks like this:

357



358

This modification ensures that F qualifies as a sheaf, since it ensures that all gluable selections of patch candidates (including the empty one) glue uniquely.

359
360

1.6. A Canonical Sheaf Construction

361

Not every presheaf is a sheaf, since some presheaves fail the gluing condition. However, there is a canonical procedure called “sheafification” that turns any presheaf into a sheaf. To sheafify a presheaf, add any missing sections that glue, then quotient sections that are locally indistinguishable. The result is guaranteed to be a sheaf, by construction.

362
363
364
365

For our purposes, there is a simplified version of sheafification that we can use to construct sheaves that model part-whole complexes in a natural way. Given a presentation of a locale, the recipe to build such a sheaf over it goes like this:

366
367
368

1. Assign local data to atomic regions.
2. Specify a gluing condition.
3. Recursively glue more and more pieces together until you can't glue any more.

369
370
371

Let's make this more precise. Given a presented locale, we can uniquely write each region as the join of its atomic regions.

372
373

Definition 14 (Atomic indices). Let $\mathbb{L} = \langle G, R \rangle$ be a presented locale with $G = \{U_1, \dots, U_n\}$. Let $\mathbb{A} \subseteq \{0, \dots, n\}$ be the indices of the atomic regions of G .

374
375

For any $U \in \mathbb{L}$, define its atomic support (denoted $I(U)$, or just I for short) as:

$$I(U) = \{i \in \mathbb{A} \mid U_i \preceq U\}$$

Then U can be written canonically as $U_{I(U)}$, the join of its atomic supports:

$$U_{I(U)} = \bigvee_{i \in I(U)} U_i.$$

Say something along the lines that there are many kinds of sheafs, and the fibers can have many different shapes. For instance, a standard example of a sheaf over a topological space is the set of real-valued functions defined over each region of that topology.

We're going to focus on a simple kind of sparse sheaf. We will build this sheaf by assigning to atomic regions some basic data: to each atomic region U_k , we assign some piece of data $\{\langle b_k \rangle\}$. What b_k is doesn't matter for our purposes at the moment. It just needs to be some piece of atomic data.

Then, we will glue together a selection of patch candidates $\{\langle b_1 \rangle\}, \{\langle b_2 \rangle\}$, and so on of atomic pieces of data to form $\{\langle b_1, b_2, \dots \rangle\}$. To decide when to glue such tuples together, we will define a gluing condition.

Talk about how we are going to imagine that each " b_k " represents something that occupies the region U_k . E.g., it's a chunk of matter, or a mechanical gadget, or whatever. But it is the part or stuff in that atomic region. The larger whole will then be built up by fusing together these parts according to the gluing condition.

A gluing condition is a family of predicates that say when a selection of patch candidates glue.

Definition 15 (Gluing condition). A gluing condition \mathcal{G} is a family of predicates

$$\mathcal{G}_U : \prod_{i \in I(U)} F(U_i) \rightarrow \{\text{true, false}\},$$

i.e. for each region $U \in \mathbb{L}$ a predicate over the Cartesian product of the fibers of the atomic support of U , that collectively satisfy the following coherence conditions:

(G1) Local data is glued. If $U_k \in \text{Atoms}(\mathbb{L})$ and $F(U_k) = \{\langle b_k \rangle\}$, then

$$\mathcal{G}_{U_k}(\langle b_k \rangle) = \text{true}.$$

(G2) Downward stability. If $\mathcal{G}_U(\langle b_i \rangle_{i \in I(U)}) = \text{true}$, then for each $V \preceq U$, we must have:

$$\mathcal{G}_V(\rho_V^U(\langle b_i \rangle_{i \in I(U)})) = \text{true}.$$

(G3) Upward stability. Given a selection of patch candidates $\langle b_i \rangle_{i \in I(U)}$, if $\mathcal{G}_{U_i \vee U_j}(\langle b_i, b_j \rangle) = \text{true}$ for each $i, j \in I(U)$, then we must have:

$$\mathcal{G}_U(\langle b_i \rangle_{i \in I(U)}) = \text{true}.$$

A sheaf can be generated from a gluing condition by starting with some local data on the atomic regions and then gluing all pieces together that satisfy the gluing condition.

Definition 16 (\mathcal{G} -sheaves). Given a gluing condition \mathcal{G} and local data $F(U_k) = \{\langle b_k \rangle\}$ for each atomic region U_k , define for each region U :

$$F(U) = \{\langle b_i \rangle_{i \in I(U)} \in \prod_{i \in I(U)} F(U_i) \mid \mathcal{G}_U(\langle b_i \rangle_{i \in I(U)}) = \text{true}\}.$$

For $V \preccurlyeq U$, define the restriction map

$$\rho_V^U : F(U) \rightarrow F(V) \quad \text{as} \quad \rho_V^U(\langle b_i \rangle_{i \in I(U)}) = \langle b_i \rangle_{i \in I(V)}.$$

Set $F(\perp) = \{\langle \rangle\}$, the empty tuple.

Remark 11. Alternatively, given some local data and a gluing condition, define a presheaf over the given locale, call it F_\wp , that assigns all combinations of local data to each region:

$$F_\wp(U) = \prod_{i \in I(U)} F(U_i).$$

Then filter by the gluing condition. That produces the same sheaf.

We must check that Definition 16 really defines a sheaf.

Theorem 1 (\mathcal{G} -sheaves are presheaves). Given a gluing condition \mathcal{G} and an assignment of local data to the atomic regions of the underlying locale, the corresponding \mathcal{G} -sheaf is a presheaf.

Proof. We must show that restrictions preserve identities and composition.

- *Identities.* ρ_U^U projects to the same index set, so $\rho_U^U = id_{F(U)}$.
- *Composition.* ρ_V^U restricts to the fiber over V , and ρ_W^V restricts to the fiber over W , so $\rho_W^V \circ \rho_V^U = \rho_W^U$.

We must also show that the restrictions are well defined.

- For $V \preccurlyeq U$, if $\langle b_i \rangle_{i \in I(U)} \in F(U)$, then by (G2) $\langle b_i \rangle_{i \in I(V)}$ satisfies \mathcal{G}_V , so ρ_V^U is well-defined. \square

Theorem 2 (\mathcal{G} -sheaves are sheaves). Given a gluing condition \mathcal{G} and an assignment of local data to the atomic regions of the underlying locale, the corresponding \mathcal{G} -sheaf satisfies the gluing condition G0.

Proof. We must show that every gluable selection of patch candidates $\langle b_i \rangle_{i \in I(U)}$ glues to yield a unique section in $F(U)$. Assume that we have a compatible selection of patch candidates $\langle b_i \rangle_{i \in I(U)}$. Then:

- Existence: we assumed the patch candidates are compatible. By (G3) then, $\langle b_i \rangle_{i \in I(U)} \in F(U)$.
- Uniqueness: let $s = \langle b_i \rangle_{i \in I(U)} \in F(U)$. If another section $t = \langle b_i \rangle_{i \in I(U)} \in F(U)$ were glued from the same components, then $s = t$, since both restrict to the same supports. \square

Throughout the rest of this paper, we will use \mathcal{G} -sheaves to model part-whole complexes, but that is only for simplicity of exposition. Any sheaf over a locale would do just as well.

2. Modeling Part-Whole Complexes as Sheaves

As noted in ??, the central claim of this paper is that we can model part-whole complexes as sheaves over locales. In particular, the locale provides the abstract parts space of “regions”

the pieces can occupy, the sheaf assigns actual pieces to those regions, and the gluing condition determines when pieces fuse.

We can thus define the core mereological concepts of part and whole in sheaf-theoretic terms. Regarding wholes, we can identify fusion with gluing: to say that some pieces fuse or form a “fusion” is just to say that they are glued together. Regarding parts, to say that a piece is a “part” is just to say that it is a part of a fusion. In other words, the parts of a fusion are just the pieces from which it is glued together.

Definition 17 (Fusions and parts). *We say that a section $s \in F(U)$ is a fusion iff*

$$\mathcal{G}_U(s) = \text{true}.$$

Given $t \in F(V)$ and $s \in F(U)$ with $V \preceq U$ and $V \neq \perp$, we say t is a part of s , denoted $t \sqsubseteq s$, iff

$$\mathcal{G}_U(s) \quad \text{and} \quad \rho_V^U(s) = t.$$

Remark 12. *$V \neq \perp$ because no parts can occupy \perp . The least region of the locale represents the combinatorial idea of no regions at all, and so it cannot be populated by any parts (hence in a \mathcal{G} -sheaf the sole section over \perp is the singleton $\langle \rangle$).*

Sheaf theory thus provides a systematic framework with which to model a large variety of part-whole complexes in a “fusions-first” manner. In the rest of this section, we illustrate with examples.

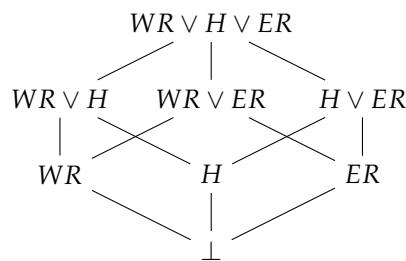
Example 14. Consider a building with a west room, an east room, and a hallway between them. For simplicity, let us consider only the floors of the building (ignore walls, ceilings, and so on). The ambient locale is given by the presentation

- $\mathbb{L} = \langle G, R \rangle = \langle \{WR, H, ER\}, \emptyset \rangle$

where

- $WR = \text{west room}$
- $H = \text{hallway}$
- $ER = \text{east room}$

As a Hasse diagram:



All of the generators are atomic, since none overlap (there are no meets among the generators):

- $\text{Atoms}(\mathbb{L}) = \{WR, H, ER\}$

Let us define a \mathcal{G} -sheaf F that models the flooring of this building. For data, let there be the following available flooring materials:

- $M = \{\text{wood, tile, ...}\}$

For a gluing condition, let us say that sections glue if they contain the same materials:

- $\mathcal{G}_U(\langle b_i \rangle_{i \in I(U)}) = \text{true iff } b_i = b_j \text{ for every } i, j \in I(U).$

- false otherwise

464

We must check that this is a legitimate gluing condition.

465

Proof. We must show that \mathcal{G} satisfies the coherence conditions (G1)–(G3).

466

G1 *Local atomic data.* Trivial.

467

G2 *Downward stability.* We must show that if $\mathcal{G}_U(\langle b_i \rangle_{i \in I(U)}) = \text{true}$, then $\mathcal{G}_V(\rho_V^U(\langle b_i \rangle_{i \in I(U)})) = \text{true}$ for every $V \preceq U$. Assume the antecedent. Then $\rho_V^U(\langle b_i \rangle_{i \in I(U)}) = \langle b_i \rangle_{i \in I(V)}$.

468

469

- Case 1: if the length of $\langle b_i \rangle_{i \in I(V)} = 1$, it glues by (G1).

470

- Case 2: if the length of $\langle b_i \rangle_{i \in I(V)} \geq 2$, then by the assumption, $b_i = b_j$ for every $i, j \in I(V)$, so they glue.

471

472

G3 *Upward stability.* Given a selection of compatible patch candidates $\langle b_i \rangle_{i \in I(U)}$, we must show that if $\mathcal{G}_{U_i}(\langle b_i, b_j \rangle) = \text{true}$ for each $i, j \in I(U)$, then $\mathcal{G}_U(\langle b_i \rangle_{i \in I(U)}) = \text{true}$. Assume the antecedent. Since for every $i, j \in I(U)$, $b_i = b_j$ by the assumption, $\langle b_i \rangle_{i \in I(U)}$ glues. \square

473

474

475

476

For the atomic regions, fix a choice of local data:

477

- $F(WR) = \{\langle \text{wood} \rangle\}$
- $F(H) = \{\langle \text{wood} \rangle\}$
- $F(ER) = \{\langle \text{tile} \rangle\}$

478

479

480

Extend compatible data to meets, of which there is only \perp , so:

481

- $F(\perp) = \{\langle \rangle\}$

482

Recursively extend data to joins via gluing:

483

- $F(WR \vee H) = \{\langle \text{wood}, \text{wood} \rangle\}$, since $F(WR) = F(H) = \{\langle \text{wood} \rangle\}$, and $\text{wood} = \text{wood}$.
- $F(WR \vee ER) = \emptyset$, since $F(WR) = \{\langle \text{wood} \rangle\}$, $F(ER) = \{\langle \text{tile} \rangle\}$, and $\text{wood} \neq \text{tile}$.
- $F(H \vee ER) = \emptyset$, since $F(H) = \{\langle \text{wood} \rangle\}$, $F(ER) = \{\langle \text{tile} \rangle\}$, $\text{wood} \neq \text{tile}$.
- $F(WR \vee H \vee ER) = \emptyset$, since $F(WR \vee H) = \{\langle \text{wood}, \text{wood} \rangle\}$, $F(H \vee ER) = F(WR \vee ER) = \emptyset$, and $\text{wood} \neq \emptyset$.

484

485

486

487

488

In this building, there are two maximal fusions:

489

- The flooring of the west room and the hallway glue into one piece that covers both.
- The flooring that covers the east room is (trivially) glued into a single piece, namely itself.

490

491

Thus, the flooring of this building is really a collection of two independent fusions: the wooden floor that covers the east room and hallway, and the tiled floor that covers the east room. That implies:

492

493

- To separate the floors of the east room and hallway, you would have to use a saw to cut them, since they are fused. They are not merely sitting next to each other. Rather, they make up a single (fused) piece.
- By contrast, to separate the hallway and the east room, you would not need to cut them, since they are not fused. They simply happen to be sitting next to each other.

494

495

496

497

498

The parts of the fusions are clear:

499

- The wooden floor that covers the west room and the hallway has two parts: the wooden floor that covers the west room, and the wooden floor that covers the hallway.
- The tiled floor of the east room has no parts (in this locale), since it is not the fusion of other fusions.

500

501

502

503

In the previous example, none of the atomic regions overlapped. The locale was discrete, and thus the sheaf was free to glue or not glue pieces as it saw fit. The story is different if there are overlaps in the locale itself. Overlaps in the locale require overlaps in the sheaf.

504

505

506

507

Example 15. Consider the floor of a single room. Let us say that the regions of interest are its west half, its east half, and a six inch span where they overlap.

The ambient locale of this kind of space can be given by the presentation

- $L = \langle G, R \rangle = \langle \{\perp, WH, O, EH\}, \{\perp \preceq O, O \preceq WH, O \preceq EH\} \rangle$

where

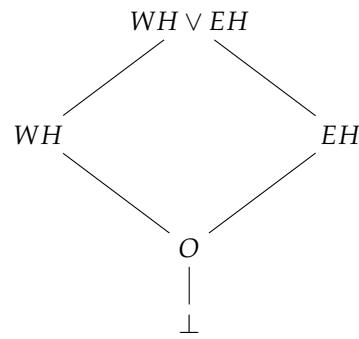
- $WH = \text{west half}$
- $O = \text{overlap}$
- $EH = \text{east half}$

The atomic sections of this locale are:

- WH
- EH
- \perp

In particular, O is not an atomic region, because it is the non-trivial overlap of WH and EH .

Here is the Hasse diagram:

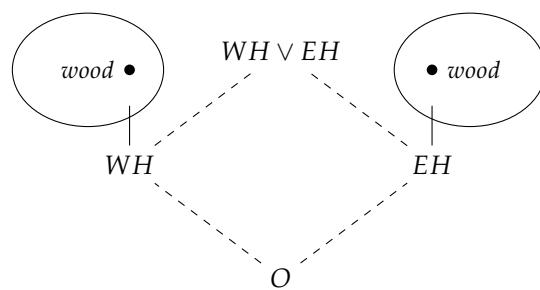


Let us define a \mathcal{G} -sheaf F that models the flooring of this room, using the same gluing condition from Example 14.

For the atomic regions, let us assign wood to both halves:

- $F(WH) = \{\langle \text{wood} \rangle\}$
- $F(EH) = \{\langle \text{wood} \rangle\}$

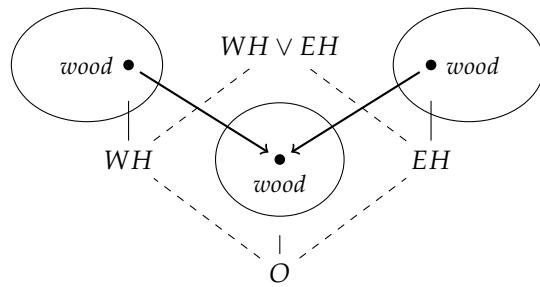
In a picture (omitting \perp for simplicity):



For the meet (the overlap O), the two halves restrict to the same thing:

- $F(O) = \{\langle \text{wood} \rangle\}$

Thus:



533

For the join, the west and east halves glue, since they're made from the same flooring materials and agree on their overlap:

534

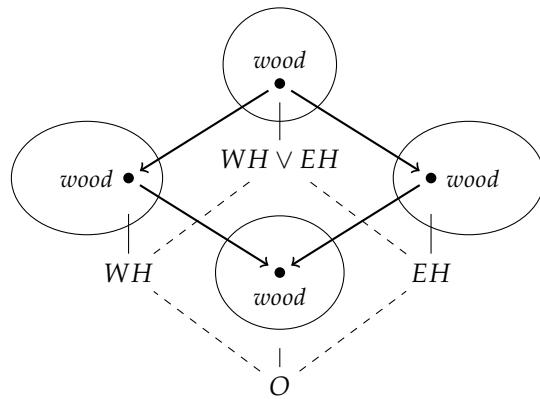
535

- $F(WH \vee EH) = \{\langle wood \rangle\}$

536

Thus:

537



538

The maximal fusion is a single piece of wooden flooring that covers the whole room. Its parts are the west and east halves, and (transitively) their overlap. The west and east halves themselves have a shared part, the strip of overlap.

539

540

541

Example 16. To illustrate a failed attempt to build a sheaf, let us take the locale and gluing condition from Example 15, but let's assign different flooring materials to the atomic regions:

542

543

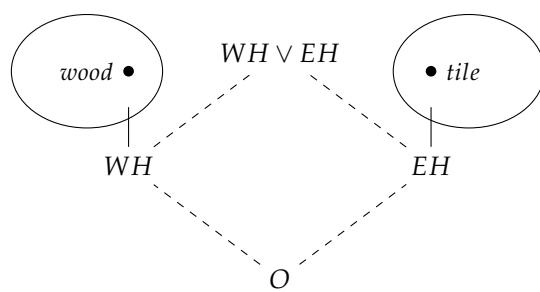
- $F(WH) = \{\langle wood \rangle\}$
- $F(EH) = \{\langle tile \rangle\}$

544

545

As a picture:

546



547

Next, we attempt to extend this data to meets, which requires that we restrict to the overlap, and then filter out anything that can't glue. Here, $F(WH)$ restricts to $\{\langle wood \rangle\}$, and $F(EH)$ restricts to $\{\langle tile \rangle\}$:

548

549

550

- $\rho_O^{WH}(\langle wood \rangle) = \langle wood \rangle$
- $\rho_O^{EH}(\langle tile \rangle) = \langle tile \rangle$

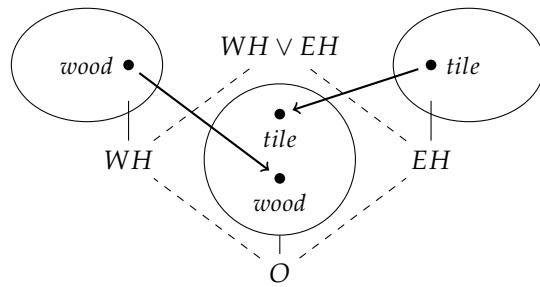
551

552

553

Thus:

553



554

However, these cannot glue, because they are not the same. We see here that the data of WH and EH disagree on the overlap. Hence, we are unable to construct a coherent sheaf. This illustrates how sheaf theory requires and manages coherent gluing at all levels. Because it requires that pieces glue together coherently at every level of “zoom,” it prevents us from ever putting together an incoherent part-whole complex in the first place.

555

556

557

558

559

It is worth spelling the failure out explicitly. Since WH and EH disagree on their overlap, F cannot assign anything to O , so:

560

561

- $F(O) = \emptyset$

562

But that renders the restrictions ρ_O^{WH} and ρ_O^{EH} undefined, thereby severing our ability to zoom in and out. Thus, the system as a whole becomes incoherent.

563

564

Intuitively, this makes sense. If the western and eastern halves of a room were truly floored with different materials, then they would not overlap. There would be some sort of boundary between them where the one’s materials end and the other’s materials begin. But here, the ambient locale doesn’t allow that possibility. In this particular locale, the western and eastern halves do overlap, so the sheaf must assign pieces to the different regions coherently, i.e., it must assign pieces that agree on their overlap.

565

566

567

568

569

570

The previous two examples were spatial. But parts come in non-spatial guises too, and sheaves can model them just as well.

571

572

Example 17. Suppose we say that human society (under some description) consists of the mesh of a specified set of relationships between the people that participate in that society.

573

574

Let P be the population in question (a finite set of individual people), and let the regions of our locale be subsets of such individuals. Then the ambient locale is given by the presentation:

575

576

- $\mathbb{L} = \langle G, R \rangle = \langle P, \emptyset \rangle$

577

For concreteness, suppose:

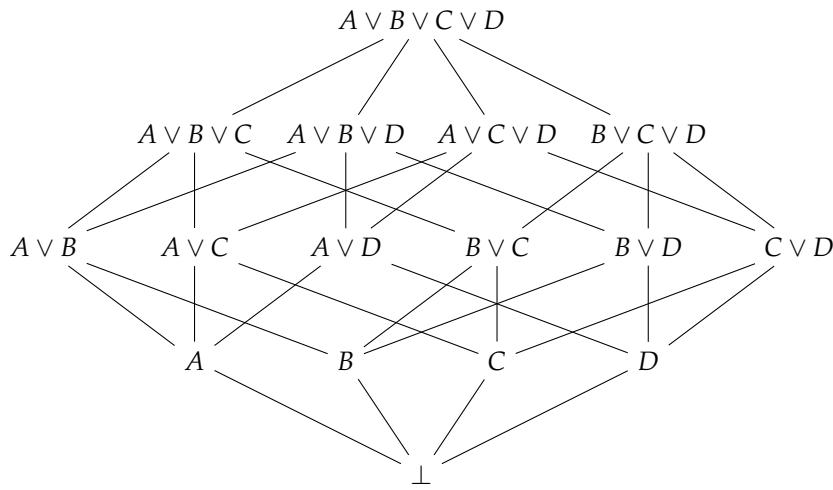
578

- $P = \{A (\text{Alice}), B (\text{Bob}), C (\text{Carol}), D (\text{Denny})\}$

579

Then the Hasse diagram is isomorphic to the powerset of P :

580



581

All of the generators are atomic, since there are no meets among the generators:

582

- Atoms(\mathbb{L}) = { A, B, C, D }

583

Let us define a \mathcal{G} -sheaf F that models the mesh of a selected set of relationships over P . To do that, let us first specify a set R that picks out the (binary, symmetric) relationships of interest:

584

585

- $R = \{f(\text{friends}), m(\text{married}), \dots\}$

586

For convenience, if $U, V \in P, r \in R$, and U and V stand in relationship r , we will write $r(U, V)$.

587

For a gluing condition, let us say that sections glue if they are connected by the same relations:

588

589

590

591

- $\mathcal{G}_U(\langle b_i \rangle_{i \in I(U)}) = \text{true iff for every } r \in R, r(U_i, U_j) \in b_i \text{ iff } r(U_j, U_i) \in b_j, \text{ for every } i, j \in I(U)$
- false otherwise

592

We must check that this is a legitimate gluing condition.

592

Proof. The proof is the same as before. \square

593

For the atomic regions, let us fix a choice of local data by assigning to each person the relations they stand in, e.g.:

594

595

- $F(A) = \{\langle \{f(A, B), f(A, C), m(A, B)\} \rangle\}$
- $F(B) = \{\langle \{f(B, A), m(B, A), f(B, D)\} \rangle\}$
- $F(C) = \{\langle \{f(C, A), m(C, D)\} \rangle\}$
- $F(D) = \{\langle \{f(D, B), m(D, C)\} \rangle\}$

596

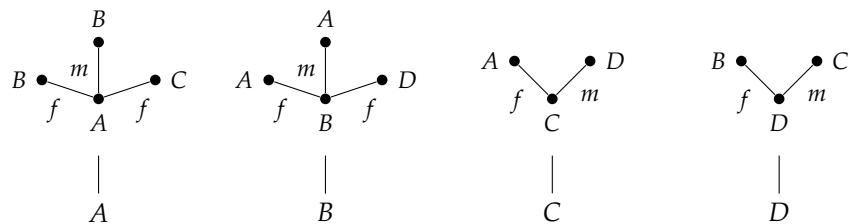
597

598

599

To visualize this data, we can picture each fiber as a mini-graph:

600



601

For example, in the fiber over A :

602

- The f -labeled edge from A to B represents $f(A, B)$: A and B are friends.
- The m -labeled edge from A to B represents $m(A, B)$: A and B are married.
- The f -labeled edge from A to C represents $f(A, C)$: A and C are friends.

603

604

605

Next, we must extend compatible data to meets, of which there is only \perp , so:

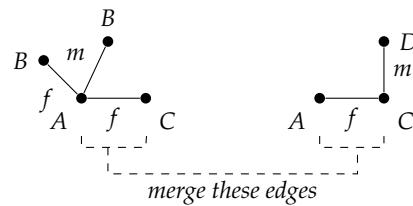
606

- $F(\perp) = \{\langle \rangle\}$

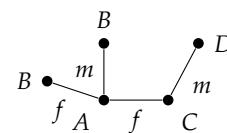
Finally, we must extend atomic data to binary joins via gluing. The gluing condition essentially says that mini-graphs can be glued along shared edges, provided that they share exactly the same edges. To see how this works, consider (for example) the mini-graphs over A and C :



Can these be glued? The answer is yes, because they share exactly one edge, namely the one labeled f . If you rotate the graphs sideways a bit, you can see how they can be merged along $f(A, C)$:



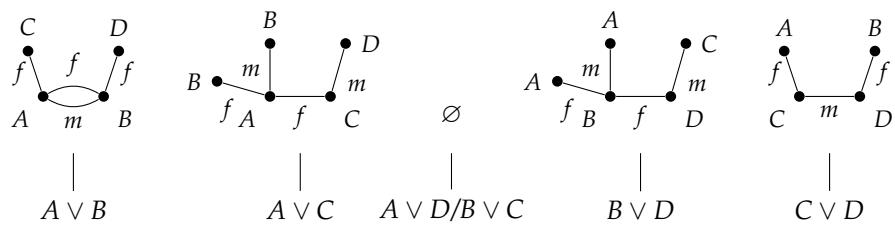
Merging along $f(A, C)$ yields the following glued graph:



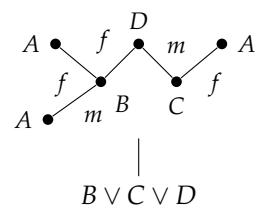
By gluing binary joins in this fashion, we get:

- $F(A \vee B) = \{\langle \{f(A, B), m(A, B), f(A, C)\}, \{f(B, A), m(B, A), f(B, D)\} \rangle\}$
- $F(A \vee C) = \{\langle \{f(B, A), m(B, A), f(B, D)\}, \{f(C, A), m(C, D)\} \rangle\}$
- $F(A \vee D) = \emptyset$
- $F(B \vee C) = \emptyset$
- $F(B \vee D) = \{\langle \{f(B, A), m(B, A), f(B, D)\}, \{f(D, B), m(D, C)\} \rangle\}$
- $F(C \vee D) = \{\langle \{f(C, A), m(C, D)\}, \{f(D, B), m(D, C)\} \rangle\}$

As pictures:



Having glued joins of two regions, we must next glue joins of three atomic regions. For instance, take $B \vee C \vee D$. We can glue $B \vee C$ trivially (because they share no edges), we can glue $C \vee D$ along their shared f -edge, and we can glue $B \vee D$ along their shared f -edge. That yields:



By gluing all joins of three atomic regions in this fashion, we get:

$$\bullet \quad F(A \vee B \vee C) = \left\{ \langle \begin{array}{l} \{f(A, B), f(A, C), m(A, B)\}, \\ \{f(B, A), m(B, A), f(B, D)\}, \\ \{f(C, A), m(C, D)\} \end{array} \rangle \right\} \quad 631$$

$$\bullet \quad F(A \vee B \vee D) = \left\{ \langle \begin{array}{l} \{f(A, B), f(A, C), m(A, B)\}, \\ \{f(B, A), m(B, A), f(B, D)\}, \\ \{f(D, B), m(D, C)\} \end{array} \rangle \right\} \quad 632$$

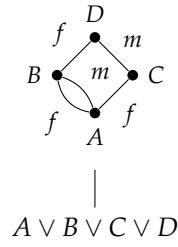
$$\bullet \quad F(A \vee C \vee D) = \left\{ \langle \begin{array}{l} \{f(A, B), f(A, C), m(A, B)\}, \\ \{f(C, A), m(C, D)\}, \\ \{f(D, B), m(D, C)\} \end{array} \rangle \right\} \quad 633$$

$$\bullet \quad F(B \vee C \vee D) = \left\{ \langle \begin{array}{l} \{f(B, A), m(B, A), f(B, D)\}, \\ \{f(C, A), m(C, D)\}, \\ \{f(D, B), m(D, C)\} \end{array} \rangle \right\} \quad 634$$

At the top-most join, gluing four regions, we get:

$$\bullet \quad F(A \vee B \vee C \vee D) = \left\{ \langle \begin{array}{l} \{f(A, B), f(A, C), m(A, B)\}, \\ \{f(B, A), m(B, A), f(B, D)\}, \\ \{f(C, A), m(C, D)\}, \\ \{f(D, B), m(D, C)\} \end{array} \rangle \right\} \quad 635$$

As a picture:



$$A \vee B \vee C \vee D \quad 638$$

The resulting sheaf yields a fused mesh of relationships over the population, which is glued together from smaller meshes over smaller subsets of the population.

- Each atomic fiber is a part of the whole (human society), and its data encodes the internal (relational) structure of that part.
- Mereological overlap is then modeled by shared relationships: two parts overlap if their relational graphs intersect coherently.
- Failure to glue (as in $F(A \vee D) = \emptyset$ and $F(B \vee C) = \emptyset$) reflects mereological separation: the atomic regions in question cannot be fused because they are not related in this mesh.

For another example, consider processes. A process (or more generally any sequence of events, states, etc.) can be seen as a part-whole complex too.

Example 18. Imagine a scenario where something can do one of two things repeatedly: at each step, it can do one thing ("option a") or another thing ("option b"), and then repeat the choice again.

To model this, fix a finite alphabet $\Sigma = \{a, b\}$, with "a" for "option a" and "b" for "option b." Then let Σ^* be the set of all finite sequences (words) over Σ , with ϵ denoting the empty sequence. For instance, the sequence aab represents the sequence of length 3 that picks "option a" first, then "option a" again, and then finally "option b."

Let us say that $\Sigma^{\leq n}$ is the set of all finite sequences less than length n , and let us say that $\Sigma^{=n}$ is the set of finite sequences of exactly length n . Hence:

- $\Sigma^{=0} = \{\epsilon\}$.
- $\Sigma^{=1} = \Sigma^{\leq 1} = \{\epsilon, a, b\}$.

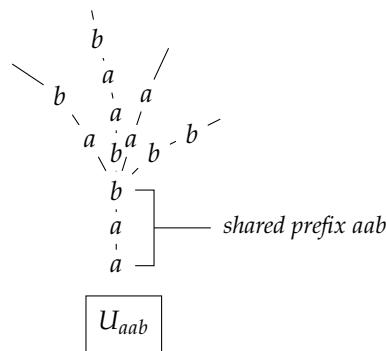
- $\Sigma^{\leq 2} = \{\epsilon, a, b, aa, bb, ab, ba\}$. 659
- $\Sigma^{=2} = \{aa, bb, ab, ba\}$. 660
- *Etc.* 661

Given sequences $w, v \in \Sigma^{\leq n}$ with $\text{length}(w) \leq \text{length}(v)$, let us write $w \subseteq v$ to denote that w is a prefix of v , as in $aab \subseteq aabc$. 662
663

Next, define a topology over $\Sigma^{\leq n}$ by setting the open sets to be sequences that share a prefix: 664

- $U_w = \{v \in \Sigma^{\leq n} \mid w \subseteq v\}$. 665

So U_w consists of all sequences that continue w . For instance, if $w = aab$, then we might picture U_w as a kind of bouquet or bundle of sequences that are all bound at their shared stem (aab) but then branch out in different directions: 666
667
668



669

We can form a locale from this topology. Let \mathbb{L} be the locale given by the presentation $\langle G, R \rangle$, where: 670

- $G = \{U_w \mid w \in \Sigma^n\}$, i.e., each open is a generator. 671
- $R = \{U_w \preceq U_v \mid v \subseteq w\}$, i.e., bouquets with longer prefixes are lower. 672

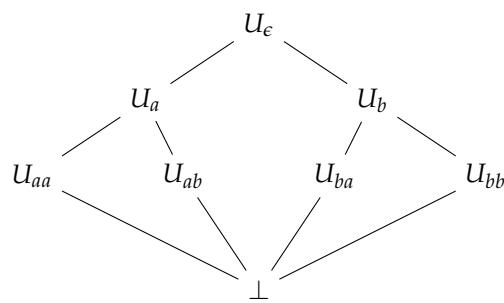
For example, given $\Sigma^{\leq 2}$, we have the following generators: 673

- $G = \{U_\epsilon, U_a, U_b, U_{aa}, U_{bb}, U_{ab}, U_{ba}\}$. 674

Here are some of the relations: 675

- $U_{aa} \preceq U_a$ and $U_{ab} \preceq U_a$, since "a" is a prefix of aa and ab. 676
- $U_{bb} \preceq U_b$ and $U_{ba} \preceq U_b$, since "b" is a prefix of bb and ba. 677
- Every generator is lower than U_ϵ , since ϵ (the empty sequence) is a prefix of every sequence. 678

The Hasse diagram looks like this: 679



680

Think of moving upwards in this locale as forgetting information about (or alternatively, as committing less to) the history of the sequence. For example, think of U_{ab} as a region where we know that "a" happened first and then "b" happened, but think of U_a as a region where we know only that "a" happened first and we don't know what happened after that. The top element is U_ϵ , which means we don't know anything about the sequence of actions. The \perp element indicates not that we know nothing, but that there is no sequence at all. 681
682
683
684
685
686

Notice that implication moves upwards: U_{ab} implies U_a because if I know (at U_{ab}) that “*a*” happened first and then “*b*” happened, then I certainly know that “*a*” happened first.

Further, no generator is the non-trivial overlap of other generators, so every generator is an atomic region:

- $\text{Atoms}(\mathbb{L}) = G$

As with any locale, we can write each region canonically as the join of its atomic regions:

- $U_{I(U)} = \bigvee_{i \in I(U)} U_i$

But here, what this means is that we can canonically write each region as the join of its “most specified” prefixes. For instance, $I(U_a) = \{aa, ab\}$, so:

- $U_a = U_{I(U_a)} = \bigvee \{U_{aa}, U_{ab}\}.$

This makes sense. Since U_a is a region where we know only that “*a*” happened first, it is the join of all maximal continuations that begin with “*a*.”

This particular locale is interesting because it models the “process space” of any 2-stage sequence that can make one of two choices at each stage. Let us now assign some actual processes to this ambient space, using a \mathcal{G} -sheaf.

Imagine a machine m that can run multiple concurrent processes, all of whom share the same memory. For simplicity, let us suppose that the machine has two registers ($R = \{r_1, r_2\}$), each of which can hold one bit (1 or 0). So, at any point in time the machine’s memory state $S : \{0, 1\} \times \{0, 1\}$ can be one of the following:

- $S = \{\langle 0, 0 \rangle, \langle 1, 0 \rangle, \langle 0, 1 \rangle, \langle 1, 1 \rangle\}$, with initial state $s_0 = \langle 0, 0 \rangle$.

We can think of the concurrent processes of interest as a selection of programs that we want to run on the machine all at the same time. In terms of behavior, let us say that each program-run reads a word from its input stream, one character at a time, and in response to each character, it takes one of the following actions A : it writes a value (1 or 0) to one of the registers, it writes (possibly distinct) values to both registers, or it does nothing and leaves the registers as they are:

- $A = \{\{r1 \mapsto v\}, \{r2 \mapsto v\}, \{r1 \mapsto v, r2 \mapsto w\}, \emptyset\}$, where $v, w \in \{0, 1\}$.

We can define a process (program trace) as a map from n -length words to n -length sequences of write actions, where we require that such maps agree on prefixes (since a process responding to *ab* and *aa* would do the same thing on the first *a*). This way, a program trace records for each input stream the sequence of write actions that result. For concreteness, here are two such traces:

- $f : \Sigma^{=2} \rightarrow A \times A$
 - $f(aa) = \langle \{r1 \mapsto 1\}, \{r1 \mapsto 0\} \rangle$
 - $f(ab) = \langle \{r1 \mapsto 1\}, \{r2 \mapsto 0\} \rangle$
 - $f(bb) = \langle \{r2 \mapsto 0\}, \{r2 \mapsto 1\} \rangle$
 - $f(ba) = \langle \{r2 \mapsto 0\}, \{r1 \mapsto 1\} \rangle$
- $g : \Sigma^{=2} \rightarrow A \times A$
 - $g(aa) = \langle \{r2 \mapsto 0\}, \{r2 \mapsto 0\} \rangle$
 - $g(ab) = \langle \{r2 \mapsto 0\}, \emptyset \rangle$
 - $g(bb) = \langle \{r1 \mapsto 1\}, \{r1 \mapsto 0\} \rangle$
 - $g(ba) = \langle \{r1 \mapsto 1\}, \{r1 \mapsto 1\} \rangle$

Let us now say that concurrent processes are compatible if they “play well” together, i.e., they share resources (memory) consistently. In particular, given two processes f and g , let us say:

- f and g are compatible at stage n if they write to different registers.
- f and g are compatible at stage n if they write the same value to the same register.

- f and g conflict at stage n if they write different values to the same register. 731

We can formalize this notion as a gluing condition that says a selection of patch candidates glue at U_w if they play well up to w . Fix a selection of programs $P = \{f, g, \dots\}$ to run on the machine, then: 732
733
734

- Given a selection of patch candidates $\langle (b_{i,p})_{p \in P} \rangle_{i \in I(U_w)}$ over a region U_w with trace length $|w|$, $\mathcal{G}_{U_w}(\langle b_i \rangle_{i \in I(U_w)}) = \text{true}$ iff the following condition holds. Write $b_{i,p}[m]$ to denote the write actions of process p in region i at stage m . Then, require that at each stage $m \leq |w|$ and every register $r \in R$, the set 735
736
737
738

$$\{v \in \{0, 1\} \mid \exists i \in I(U_w), p \in P \text{ where } r \mapsto v \in b_{i,p}[m]\}$$

has cardinality at most 1. In other words, two values are not written to the same register. 739

- $\mathcal{G}_{U_w}(\langle b_i \rangle_{i \in I(U_w)}) = \text{false}$ otherwise. 740

We must check that this is a legitimate gluing condition. 741

Proof. We must check that \mathcal{G}_{U_w} is downwards and upwards stable. 742

- *Downwards stability.* Assume that $\mathcal{G}_{U_w}(\langle b_i \rangle_{i \in I(U_w)}) = \text{true}$. Then $\mathcal{G}_{U_i}(\langle b_i \rangle) = \text{true}$ for each $i \in I(U_w)$ since by \mathcal{G}_{U_w} , every b_i, b_j play well on their prefixes. 743
744
- *Upwards stability.* Assume $\mathcal{G}_{U_{\{i,j\}}}(\langle b_i, b_j \rangle) = \text{true}$ for all $i, j \in I(U_w)$. Then $\mathcal{G}_{U_w}(\langle b_i \rangle_{i \in I(U_w)}) = \text{true}$ since no i, j conflict on writes. \square 745
746

With a gluing condition at hand, we can now define a \mathcal{G} -sheaf F . Let our selection of processes be $P = \{f, g, \dots\}$. Then, we can fix the atomic data (omitting outer brackets to avoid clutter): 747
748

- $F(U_{aa}) = (f(aa), g(aa)) = (\langle \{r1 \mapsto 1\}, \{r1 \mapsto 0\} \rangle, \langle \{r2 \mapsto 0\}, \{r2 \mapsto 0\} \rangle)$. 749
- $F(U_{ab}) = (f(ab), g(ab)) = (\langle \{r1 \mapsto 1\}, \{r2 \mapsto 0\} \rangle, \langle \{r2 \mapsto 0\}, \emptyset \rangle)$. 750
- $F(U_{bb}) = (f(bb), g(bb)) = (\langle \{r2 \mapsto 0\}, \{r2 \mapsto 1\} \rangle, \langle \{r1 \mapsto 1\}, \{r1 \mapsto 0\} \rangle)$. 751
- $F(U_{ba}) = (f(ba), g(ba)) = (\langle \{r2 \mapsto 0\}, \{r1 \mapsto 1\} \rangle, \langle \{r1 \mapsto 1\}, \{r1 \mapsto 1\} \rangle)$. 752

There are no meets among the generators beyond \perp , so: 753

- $F(\perp) = \{\langle \rangle\}$. 754

Next, we must extend F to joins via gluing. So, for each U_w , we must assign: 755

- $F(U_w) = \{\langle (b_{i,p})_{p \in P} \rangle_{i \in I(U_w)} \mid \mathcal{G}_{U_w}(\langle (b_{i,p})_{p \in P} \rangle_{i \in I(U_w)}) = \text{true}\}$. 756

Let's compute $F(U_a) = F(U_{aa} \vee U_{ab})$. To determine if $(f(aa), g(aa))$ and $(f(ab), g(ab))$ glue, we need to check that they do not write conflicting values. 757
758

- Stage 1 (at the shared prefix "a"): $f(aa)$ and $f(ab)$ write 1 to $r1$, while $g(aa)$ and $g(ab)$ write 0 to $r2$. Since f and g write to different registers, there is no conflict. 759
760
- Stage 2: $f(ab)$ and $g(aa)$ write the same value (namely, 0) to $r2$, $f(aa)$ writes 0 to $r1$, and $g(ab)$ does nothing, so there are no conflicts. 761
762

Hence, $(f(aa), g(aa))$ and $(f(ab), g(ab))$ glue to form a unique section: 763

- $F(U_a) = (f(aa), g(aa)), (f(ab), g(ab))$ 764

Now let's compute $F(U_b) = F(U_{bb} \vee U_{ba})$. To determine if $(f(bb), g(bb))$ and $(f(ba), g(ba))$ glue, we need to again check for conflicting writes: 765
766

- Stage 1 (at the shared prefix "b"): $f(bb)$ and $f(ba)$ write 0 to $r2$, while $g(bb)$ and $g(ba)$ write 1 to $r1$, so there is no conflict. 767
768
- Stage 2: $f(bb)$ and $f(ba)$ write 1 to different registers, so they do not conflict with each other, while $f(ba)$ and $g(ba)$ write 1 to $r1$, so they do not conflict either. However, $g(bb)$ writes 0 to $r1$, which conflicts with $g(ba)$'s and $f(ba)$'s attempt to write 1 to the same register. 769
770
771

Since we have a conflict, $(f(bb), g(bb))$ and $(f(ba), g(ba))$ fail to glue over U_b . Notice:

- The processes f and g agree locally at U_a .
- By contrast, they disagree locally at U_b .
- There is no global section that glues together all of f and g 's behavior at the top U_e , thus f and g are not globally compatible processes.

This sort of example illustrates how sheaves can model processes, concurrency, and resource conflicts. Here the processes were programs running on a simple machine, but they could just as easily be biological processes competing for resources, etc.

Whatever the concrete details may be, this example captures how local behavior can be integrated and extended over larger regions of the process space. One might naively think that the "parts" of such systems are the processes. But there is a different way to slice it: if you want to talk about the integrity of the "whole" of a concurrent system, you need to talk about how that involves coherent, integrated behavior that is functionally united locally across the various "regions" and "stages" of the system's evolution.

TODO:

- Add example: something over a continuous interval/timeline? E.g., maybe something over a timeline (the frame of opens taken from the usual topology of \mathbb{R})? Maybe we can define a gluing condition for a mass of clay over time that says pieces glue if they agree on overlaps, so that the whole lump of clay can have parts replaced over time but as a whole it never breaks into fragments? Maybe the "closure" is even a modality.
- Add example: Socrates and seated Socrates?
- Note Spivak et al's behavioral mereology is an example of a \mathcal{G} -sheaf (and check the details to make sure that's really true).
- Mormann's "structural mereology" is basically just our thesis. Add examples from his similarity structures.

3. Modalities in the Sheaf-theoretic Setting

In the context of sheaves, modalities manifest as j -operators (also called local operators). A j -operator is a closure operator on the underlying locale.

Definition 18 (j -operators). Given a locale \mathbb{L} , a j -operator on \mathbb{L} is a closure operator $j : \mathbb{L} \rightarrow \mathbb{L}$ satisfying the following conditions:

- (J1) Inflation. $U \preceq j(U)$.
- (J2) Idempotence. $j(j(U)) = j(U)$.
- (J3) Meet-preservation. $j(U \wedge V) = j(U) \wedge j(V)$.

A j -operator induces a j -sheaf.

Definition 19 (j -sheaves). Given a sheaf F over a locale \mathbb{L} and a j -operator $j : \mathbb{L} \rightarrow \mathbb{L}$, the corresponding j -sheaf, denoted F_j , is given by:

$$F_j = F(j(U)).$$

Remark 13. In a sheaf, there are a variety of other modalities beyond the traditional alethic ones (necessity and possibility). Any closure operator qualifies as a modality of some description.

Example 19. From Example 17, recall the mesh of human relationships modeled by a \mathcal{G} -sheaf F defined over the presented locale $\mathbb{L} = \langle P := \{A, B, C, D\}, \emptyset \rangle$. Let us define a family of "reachability" modalities over this mesh.

For each $r \in R$, write \rightsquigarrow_r for the reflexive and transitive closure of r on the generators. Hence, \rightsquigarrow_f is the transitive closure of friendship on the generators, and \rightsquigarrow_m is the transitive closure on marriage.

Then for each $r \in R$, define j_r inductively:

- Base case. For each generator $U \in G$, set j_r to the join of all other generators reachable via r :

$$j_r(U) := \bigvee \{V \mid U \rightsquigarrow_r V\}$$

- Inductive step. Extend to arbitrary joins $U_{I(U)}$:

$$j_r(U_{I(U)}) := \bigvee_{i \in I(U)} j_r(U_i)$$

We need to check that this is a j -operator.

Proof. We check (J1)–(J3) from the definition.

Do the base case, then the inductive step.

□

Intuitively, this operator expands every region U to the largest region that is reachable from U by r . In other words, it expands each subset of society to the largest subset of society that is connected by r . Hence, $j_f(U)$ yields all those who are connected to U through a chain of friends, while $j_m(U)$ yields all those who are connected to U through a chain of marriage (which in a monogamous society will yield only married couples but in a polygamous society may be more revealing).

Applying j_f (for instance) to \mathbb{L} yields the following:

- $j_f(A) = A \vee B \vee C \vee D$, because $A \rightsquigarrow_f A$, $A \rightsquigarrow_f B$, $A \rightsquigarrow_f D$, and $A \rightsquigarrow_f C$.
- $j_f(B) = A \vee B \vee C \vee D$, because $B \rightsquigarrow_f B$, $B \rightsquigarrow_f D$, $B \rightsquigarrow_f A$, and $B \rightsquigarrow_f C$.
- Similar for $j_f(C)$ ad $j_f(D)$.
- $j_f(\perp) = \perp$.

Hence, everyone in this mini-society is connected through friends (or friends-of-friends, etc.). Notice also that everyone is connected immediately, i.e., at the first application of j_f .

When it comes to marriage, the situation is different. Applying j_m yields:

- $j_m(A) = A \vee B$, because $A \rightsquigarrow_m A$ and $A \rightsquigarrow_m B$.
- $j_m(B) = A \vee B$, because $B \rightsquigarrow_m B$ and $B \rightsquigarrow_m A$.
- $j_m(C) = C \vee D$, because $C \rightsquigarrow_m C$ and $C \rightsquigarrow_m D$.
- Similar for $j_m(D)$.
- $j_m(A \vee B) = A \vee B$, since A and B are already connected.
- $j_m(C \vee D) = C \vee D$, since C and D are already connected.
- $j_m(A \vee C) = A \vee B \vee C \vee D$, since from A , A can reach B (i.e., $A \rightsquigarrow_m B$) and from C , C can reach D (i.e., $C \rightsquigarrow_m D$).
- Similar for the rest.

In contrast to the j_f modality, the j_m modality keeps the A, B component separate from the C, D component at all regions (sub-populations) that don't include a member of both couples, just as we would expect.

Now that we have defined j_f and j_m , we can construct a modal overlay for each that we can use to filter the original mesh:

- Define the friendship mesh as F_f , filtered by j_f , i.e., set $F_f(U) := F(j_f(U))$.
- Define the marriage mesh as F_m , filtered by j_m , i.e., $F_m(U) := F(j_m(U))$.

Example 20. Recall the example of concurrent processes f and g from Example 18. We can define an “already happened” modality that captures what has definitely occurred so far.

Definition 20 (Already-happened operator). Let j_H be given by:

$$j_H(U_w) := \bigvee \{U_v \mid v \subseteq w\},$$

i.e., the join of all opens corresponding to prefixes of w (including w itself).

Intuitively, $j_H(U_w)$ is the region that remembers everything that has already happened along w . It is a closure operator that closes upwards by collecting all shorter prefixes.

We must check that this is a legitimate j -operator.

Proof. We check (J1)–(J3).

- J1 *Inflation.* $U_w \preceq j_H(U_w)$ holds because U_w is included among the prefixes being joined.
- J2 *Idempotence.* Applying j_H again adds no new prefixes, so $j_H(j_H(U_w)) = j_H(U_w)$.
- J3 *Meet-preservation.* The meet of two regions corresponds to their longest shared prefix, whose prefixes are all of the prefixes collected by j_H . Hence, $j_H(U_w \wedge U_v) = j_H(U_w) \wedge j_H(U_v)$. \square

Applying j_H to the generators of \mathbb{L} :

- For U_{aa} : $j_H(U_{aa}) = U_\epsilon \vee U_a \vee U_{aa}$.
- For U_{ab} : $j_H(U_{ab}) = U_\epsilon \vee U_a \vee U_{ab}$.
- For U_{ba} : $j_H(U_{ba}) = U_\epsilon \vee U_b \vee U_{ba}$.
- For U_{bb} : $j_H(U_{bb}) = U_\epsilon \vee U_b \vee U_{bb}$.
- For U_a : $j_H(U_a) = U_\epsilon \vee U_a$.
- For U_b : $j_H(U_b) = U_\epsilon \vee U_b$.
- For U_ϵ : $j_H(U_\epsilon) = U_\epsilon$.

Since j_H filters each region to everything that is already determined in that region, we can use it to define an overlay of F

$$F_H(U_w) := F(j_H(U_w)),$$

so that sections at U_w remember only what has happened along all prefixes of w .

Example 21. Recall the example of concurrent processes f and g from Example 18. We can define a safety (“nothing bad happens”) modality as a j -operator that identifies the largest safe extensions of a given region.

Definition 21 (Safety operator). Let us say that a region U is safe if all processes in $F(U)$ play well together, i.e., if there are no write conflicts. Then let $j_S : \mathbb{L} \rightarrow \mathbb{L}$ be given by:

$$j_S(U) := \begin{cases} \bigvee \{V \mid U \preceq V \text{ and } V \text{ is safe}\} & \text{if this join is non-empty} \\ U & \text{otherwise.} \end{cases}$$

Intuitively, $j_S(U)$ inflates U to the largest region that is guaranteed safe starting from U .

We must check that $j_S(U)$ is a legitimate j -operator.

Proof. We check (J1)–(J3).

- J1 *Inflation.* By construction, $U \preceq j_S(U)$ whenever U has any safe parent regions, otherwise $j_S(U) = U$.

- J2 *Idempotence.* Applying j_S more than once does not change the result, since applying it once takes the join of all safe parents. Hence, $j_S(j_S(U)) = j_S(U)$. 886
J3 *Meet-preservation.* For any U and V , since $U \wedge V$ is U or V , 887
888

$$j_S(U \wedge V) = \bigvee \{W \mid U \wedge V \preceq W \text{ and } W \text{ safe}\} = j_S(U) \wedge j_S(V). \quad \square$$

Let's apply j_S to the generators of \mathbb{L} : 889

- $j_S(U_{aa}) = U_a$ since its safe parent regions are U_{aa} and U_a . 890
- Similarly, $j_S(U_{ab}) = U_a$. 891
- $j_S(U_{ba}) = U_{ba}$ because the only safe parent of U_{ba} is U_{ba} itself. 892
- Similarly, $j_S(U_{bb}) = U_{bb}$. 893

Now extend it to joins: 894

- $j_S(U_a) = j_S(U_{aa}) \vee j_S(U_{ab}) = U_a \vee U_a = U_a$. 895
- $j_S(U_b) = U_b$ since U_b is unsafe (there are conflicts among its generators) and thus no further extension can be safe. 896
897
- \perp is trivially fixed: $j_S(\perp) = \perp$. 898

Notice: 899

- Each generator U_w represents a part of a process's history. 900
- The operator j_S identifies the largest safe fusion containing U_w , i.e., the maximal extension of the part where processes play well together. 901
902
- If no safe extensions exist (as in U_b), then $j_S(U_b)$ doesn't get bigger, indicating that safety cannot be guaranteed any further beyond this part. 903
904
- Hence, j_S captures a mereological notion of integrity, showing which combinations of parts form consistent wholes and which do not. 905
906

TODOs: 907

- Add example: A statue and the lump of clay? 908

4. Classical Mereological Notions in the Sheaf-theoretic Setting 909

In this section, we provide a discussion of what classical notions of mereology look like in the sheaf-theoretic setting. 910
911

- *Cambridge fusions.* Sheaves handle Cambridge fusions correctly. 912
- *Mere collections.* The collection of all dogs. Is that a "whole"? Well, we could build a sheaf whose atomic regions are filled with dogs, none of which glue. Then we have a collection of dogs, but no glued object. That matches exactly the intuition: yes, we have a "collection" (we built a sheaf for it, after all), but the internals of that sheaf reveal that it's merely a collection, i.e., that its parts are not glued. 913
914
915
916
917
- *Co-habitating fusions.* Sheaves allow multiple fusions to occupy the same locale, without being glued. For instance, in the sheaf of real-valued functions over real number line, there are many functions that glue together, and occupy the same locale. 918
919
920
- *Non-boolean algebra.* The parts space is Heyting, not Boolean. We're not saddled with such a strong complement operation. You can pick a locale that is Boolean if you need it, but this framework doesn't require it. In fact, the positive logic of a locale is "geometric logic." 921
922
923
924
- *Reflexivity, antisymmetry, and transitivity.* These are guaranteed. Locally, of course, you may not have transitivity. But globally, it's a theorem. [Check that.] 925
926
- *Distributivity.* 927

do the glued sections of a sheaf have to be distributive? Only inside what glues (since we glue pairwise, so every $i \vee j$ of the cover.

- *An empty element.* There is a need for a bottom element in the *algebra* of parts, but a sheaf need not contain any such thing. There is no need here to try and construct awkward mathematical structures that do algebra on parts but yet don't have a bottom element because our ontological intuitions tell us there can be no such thing. That confuses two issues: algebra and integrity. So here we separate those cleanly, and the algebra can do algebra while the sheaf can do integrity. [In a sheaf you CAN'T assign an empty element to bottom, for coherence, so the bottom element is special...need to say more about that and figure it out.]
- *Supplementation principles.* Sheaves don't constrain one way or another. [Is that really true? Maybe it's better to say that it doesn't force any supplementation principles, which might provide a reason to call into question whether supplementation is another one of those ideas that is about integrity of parts but has been confused with the algebra of parts.]
- *Ordering of parts.* Consider that "pit" and "tip" have the same parts but are different words. These differences can be handled by different sheaves over a 3-stage prefix-ordered locale as in the example of concurrent processes. Note that we retain extensionality.
- *Extensionality.* Classical mereology's notion of extensionality essentially flattens any structure and is thus overly aggressive. This is why extensionality is so controversial. The sheaf-theoretic perspective retains extensionality, but is much more nuanced. [Here too, I suspect that mereological discussions of extensionality have confused the algebra of parts and the integrity of wholes.]
- *Gunk and atoms.* You can model continuity and gunky parts if you so desire. You just need a sober space to do it.

check that we can model continuity in the locale in this way.

can you do continuity only in the sheaf data, without an underlying continuous decompositon in the locale? I would think that if you can't infinitely decompose into smaller opens around a point in the locale, you couldn't do such a thing in the sheaf data?

- *Priority of wholes.* The framework is agnostic as to whether you take an Aristotelian-Thomistic approach
 cite Aquinas, Arlig, and that guy who wrote that recent book defending the Aristotelian view
- *The whole is greater than its parts.* The framework is agnostic as to whether you want to be a Scotist and say that the whole is something over and above its parts (cite Cross) or an Ockhamist who says it is not
 cite Normore, Arlig

Todo list

use “primitive” or “fundamental” instead of “atomic”? An atomic thing can’t be broken down, but a meet-irreducible element does have a part!	967
7	968
Say something along the lines that there are many kinds of sheafs, and the fibers can have many different shapes. For instance, a standard example of a sheaf over a topological space is the set of real-valued functions defined over each region of that topology.	969
	970
	971
	972
	973
Talk about how we are going to imagine that each “ b_k ” represents something that occupies the region U_k . E.g., it’s a chunk of matter, or a mechanical gadget, or whatever. But it is the part or stuff in that atomic region. The larger whole will then be built up by fusing together these parts according to the gluing condition.	974
21	974
Do the base case, then the inductive step.	975
35	975
do the glued sections of a sheaf have to be distributive? Only inside what glues (since we glue pairwise, so every $i \vee j$ of the cover)	976
37	976
check that we can model continuity in the locale in this way.	977
38	977
can you do continuity only in the sheaf data, without an underlying continuous decompositon in the locale? I would think that if you can’t infinitely decompose into smaller opens around a point in the locale, you couldn’t do such a thing in the sheaf data?	978
	979
	980
	981
	982
	983
	984
cite Aquinas, Arlig, and that guy who wrote that recent book defending the Aristotelian view	985
38	985
cite Normore, Arlig	986
38	986
	987
	988